

Applied Data Analysis for Public Policy Studies

Introduction

Michele Fioretti
Sciences Po Paris
2020-08-25

Welcome to Applied Data Analysis for Public Policy Studies



Welcome to Applied Data Analysis for Public Policy Studies

- This course will teach you the core tools of *econometrics*.



Welcome to Applied Data Analysis for Public Policy Studies

- This course will teach you the core tools of *econometrics*.
- We will meet every Wednesday at XXX on Zoom.



Welcome to Applied Data Analysis for Public Policy Studies

- This course will teach you the core tools of *econometrics*.
- We will meet every Wednesday at XXX on Zoom.
- You will meet with the TA, Florin Cogne, every couple of weeks to discuss the material and review problem sets.

Welcome to ScPoEconometrics!



Welcome to Applied Data Analysis for Public Policy Studies

- This course will teach you the core tools of *econometrics*.
- We will meet every Wednesday at XXX on Zoom.
- You will meet with the TA, Florin Cogne, every couple of weeks to discuss the material and review problem sets.

Welcome to ScPoEconometrics!

- We will learn these tools through example based on *the programming language R*.



Welcome to Applied Data Analysis for Public Policy Studies

- This course will teach you the core tools of *econometrics*.
- We will meet every Wednesday at XXX on Zoom.
- You will meet with the TA, Florin Cogne, every couple of weeks to discuss the material and review problem sets.

Welcome to ScPoEconometrics!

- We will learn these tools through example based on *the programming language R*.
- This course builds on the (amazing) *openSource lectures* by the ScPo Econometrics team (Florian Oswald, Gustave Kenedi and Pierre Villedieu).



Welcome to Applied Data Analysis for Public Policy Studies

- This course will teach you the core tools of *econometrics*.
- We will meet every Wednesday at XXX on Zoom.
- You will meet with the TA, Florin Cogne, every couple of weeks to discuss the material and review problem sets.

Welcome to ScPoEconometrics!

- We will learn these tools through example based on *the programming language R*.
- This course builds on the (amazing) *openSource lectures* by the ScPo Econometrics team (Florian Oswald, Gustave Kenedi and Pierre Villedieu).
- A *book overview* of the material is also available.



Welcome to Applied Data Analysis for Public Policy Studies

- This course will teach you the core tools of *econometrics*.
- We will meet every Wednesday at XXX on Zoom.
- You will meet with the TA, Florin Cogne, every couple of weeks to discuss the material and review problem sets.

Welcome to ScPoEconometrics!

- We will learn these tools through example based on *the programming language R*.
- This course builds on the (amazing) *openSource lectures* by the ScPo Econometrics team (Florian Oswald, Gustave Kenedi and Pierre Villedieu).
- A *book overview* of the material is also available.
- The goal of the course is to provide you with essential data analysis tools (a.k.a. *econometrics*).



Answering Important Questions with Econometrics

Does immigration *cause* lower wages and higher unemployment for locals?



Answering Important Questions with Econometrics

Does immigration *cause* lower wages and higher unemployment for locals?

Does increasing the minimum wage *cause* greater unemployment?



Answering Important Questions with Econometrics

Does immigration *cause* lower wages and higher unemployment for locals?

Does increasing the minimum wage *cause* greater unemployment?

Does more education *cause* higher wages?



Answering Important Questions with Econometrics

Does immigration *cause* lower wages and higher unemployment for locals?

Does increasing the minimum wage *cause* greater unemployment?

Does more education *cause* higher wages?

Does higher public debt levels *cause* lower economic growth?



Answering Important Questions with Econometrics

Does immigration *cause* lower wages and higher unemployment for locals?

Does increasing the minimum wage *cause* greater unemployment?

Does more education *cause* higher wages?

Does higher public debt levels *cause* lower economic growth?


Does birth order *cause* differing education trajectories?



Causality

- Notice the keyword **cause** in all of the above.
- Notice also that *many other factors could have caused* each of those outcomes.
- Econometrics is often about spelling out conditions under which we can claim to measure causal relationships.
- We will encounter the most basic of those conditions, and talk about some potential pitfalls.

As in the acclaimed **Book of Why** we often ask *why* did something happen?



This Course

- Teach you the basics of *linear regression*, *statistical inference* and *impact evaluation*.



This Course

- Teach you the basics of *linear regression*, *statistical inference* and *impact evaluation*.
- Equip you with a framework to think more deeply about causality.



This Course

- Teach you the basics of *linear regression*, *statistical inference* and *impact evaluation*.
- Equip you with a framework to think more deeply about causality.
- Introduce you to the R software environment.



This Course

- Teach you the basics of *linear regression*, *statistical inference* and *impact evaluation*.
- Equip you with a framework to think more deeply about causality.
- Introduce you to the R software environment.
- ⚠ This is *not* a course about R.



This Course

- Teach you the basics of *linear regression*, *statistical inference* and *impact evaluation*.
- Equip you with a framework to think more deeply about causality.
- Introduce you to the R software environment.
- ⚠ This is *not* a course about R.

Grading

1. There will be *periodic quizzes* on Moodle roughly every two weeks => 10%



This Course

- Teach you the basics of *linear regression*, *statistical inference* and *impact evaluation*.
- Equip you with a framework to think more deeply about causality.
- Introduce you to the R software environment.
- ⚠ This is *not* a course about R.

Grading

1. There will be *periodic quizzes* on Moodle roughly every two weeks => 10%
2. There will be *two problem sets / case studies* => 40%



This Course

- Teach you the basics of *linear regression*, *statistical inference* and *impact evaluation*.
- Equip you with a framework to think more deeply about causality.
- Introduce you to the R software environment.
- ⚠ This is *not* a course about R.

Grading

1. There will be *periodic quizzes* on Moodle roughly every two weeks => 10%
2. There will be *two problem sets / case studies* => 40%
3. There will be *a take home exam* => 50%



Course Materials

1. The **Book**
2. The **Slides**
3. The code repository for the **R package**
4. Quizzes on **Moodle**



Syllabus 🙌

Lecture 1: **Introduction**

Quiz 1 (after lecture 2)

Lecture 2/3: **Summarising Data**

Quiz 2

Lecture 4: **Simple Linear Regression**

Lecture 5: **Introduction to Causality**

Midterm Project

Lecture 6: **Multiple Linear Regression**

Lecture 7: **Sampling**

Quiz 3

Lecture 8/9: **Statistical Inference**

Quiz 4

Lecture 10: **Differences-in-Differences**

Lecture 11: **Regression Discontinuity**

Quiz 5

Lecture 12: **Recap**

Final Project



Useful Resources (Other Than our *Book*)

Econometrics

- *Mastering Metrics* by Angrist and Pischke
- *Modern Introduction to Econometrics* by Wooldridge
- *Introduction to Econometrics* by Stock and Watson
- *Causal Inference: The Mixtape* by Cunningham
- Ben Lambert's youtube channel

Metrics and R

- ModernDive
- Introduction to Econometrics with R
- R for Data Science



R

What is R?

R is a **programming language** with powerful statistical and graphic capabilities.



What is R?

R is a **programming language** with powerful statistical and graphic capabilities.

Why are we using R?¹



What is R?

R is a **programming language** with powerful statistical and graphic capabilities.

Why are we using R?¹

1. R is **free** and **open source**—saving both you and the university 💰💰💰.



What is R?

R is a **programming language** with powerful statistical and graphic capabilities.

Why are we using R?¹

1. R is **free** and **open source**—saving both you and the university 💰💰💰.
2. R is very **flexible and powerful**—adaptable to nearly any task, *e.g.*, data cleaning, data visualization, econometrics, spatial data analysis, machine learning, web scraping, ...



What is R?

R is a **programming language** with powerful statistical and graphic capabilities.

Why are we using R?¹

1. R is **free** and **open source**—saving both you and the university 💰💰💰.
2. R is very **flexible and powerful**—adaptable to nearly any task, *e.g.*, data cleaning, data visualization, econometrics, spatial data analysis, machine learning, web scraping, ...
3. R has a vibrant, thriving online community that will (almost) always have a solution to your problem. (**stack overflow**)



What is R?

R is a **programming language** with powerful statistical and graphic capabilities.

Why are we using R?¹

1. R is **free** and **open source**—saving both you and the university 💰💰💰.
2. R is very **flexible and powerful**—adaptable to nearly any task, *e.g.*, data cleaning, data visualization, econometrics, spatial data analysis, machine learning, web scraping, ...
3. R has a vibrant, thriving online community that will (almost) always have a solution to your problem. (**stack overflow**)
4. If you put in the work², you will come away with a **very valuable and marketable** tool.

[1]: This list has been inspired by **Ed Rubin's**.

[2]: Learning R definitely requires time and effort but it's worth it, trust me! 💪.



Why can't we just use Excel?

Many reasons but here are just a few:



Why can't we just use Excel?

Many reasons but here are just a few:

- Not reproducible.



Why can't we just use Excel?

Many reasons but here are just a few:

- Not reproducible.
- Not straightforward to merge datasets together.



Why can't we just use Excel?

Many reasons but here are just a few:

- Not reproducible.
- Not straightforward to merge datasets together.
- Very fastidious to clean data.



Why can't we just use Excel?

Many reasons but here are just a few:

- Not reproducible.
- Not straightforward to merge datasets together.
- Very fastidious to clean data.
- Limited to small datasets



Why can't we just use Excel?

Many reasons but here are just a few:

- Not reproducible.
- Not straightforward to merge datasets together.
- Very fastidious to clean data.
- Limited to small datasets
- Not designed for proper econometric analyses, maps, complex visualisations, etc.



R SHOWCASE

Showcase #1: Spatial Data

- R is very strong with spatial data. In particular via the sf package.
- We can represent *any* shape or geometry.
- Maps are the most obvious example:

```
library(sf)
library(tmap)
iris_shfl <- read_sf("chapter1_files/figure-html/cont")
  mutate(dep = substr(INSEE_COM,1,2)) %>%
  select(CODE_IRIS, dep, geometry) %>%
  filter(dep == "75")
iris_income <- readRDS("../rds/iris_inc.rds") %>%
  mutate(CODE_IRIS = IRIS) %>%
  select(CODE_IRIS, DISP_MED15)
iris_map <- left_join(iris_shfl, iris_income, by = "C")
tmap_mode("plot")
tm_shape(iris_map) +
  tm_borders() +
  tm_fill(col = "DISP_MED15", title = "Median househo")
```

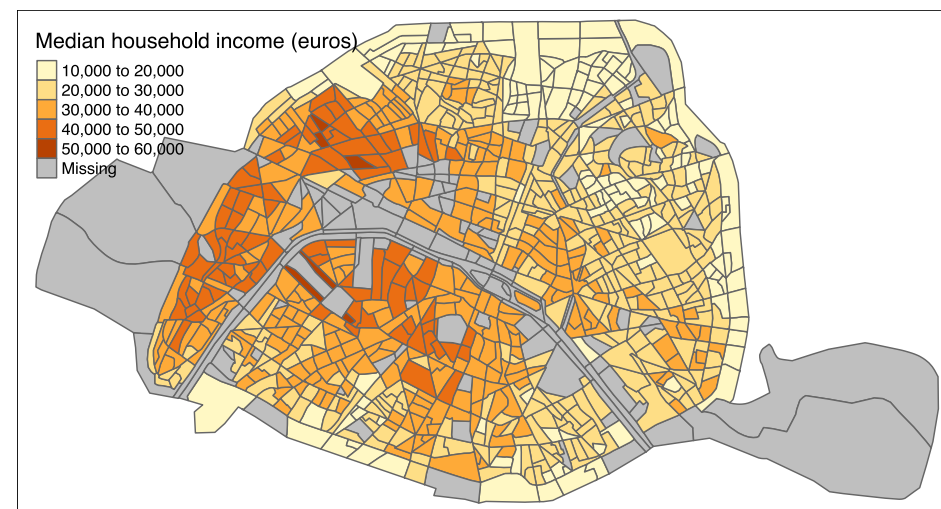


Showcase #1: Spatial Data

- **R** is very strong with spatial data. In particular via the **sf** package.
- We can represent *any* shape or geometry.
- Maps are the most obvious example:

```
library(sf)
library(tmap)
iris_shfl <- read_sf("chapter1_files/figure-html/cont")
  mutate(dep = substr(INSEE_COM,1,2)) %>%
  select(CODE_IRIS, dep, geometry) %>%
  filter(dep == "75")
iris_income <- readRDS("../rds/iris_inc.rds") %>%
  mutate(CODE_IRIS = IRIS) %>%
  select(CODE_IRIS, DISP_MED15)
iris_map <- left_join(iris_shfl, iris_income, by = "C")
tmap_mode("plot")
tm_shape(iris_map) +
  tm_borders() +
  tm_fill(col = "DISP_MED15", title = "Median househo
```

- Can be improved but you get this with only **14 lines of code!**




Showcase #2: Interactive web applications



Showcase #2: Interactive web applications

- Santé publique France has created a **simple web application** to track the epidemiological situation in French regions per week



In Practice: Data Wrangling



In Practice: Data Wrangling

- You will spend a lot of time preparing data for further analysis.



In Practice: Data Wrangling

- You will spend a lot of time preparing data for further analysis.
- The `gapminder` dataset contains data on life expectancy, GDP per capita and population by country between 1952 and 2007.
- Suppose we want to know the average life expectancy and average GDP per capita for each continent in each year.
- We need to group the data by continent *and* year, then compute the average life expectancy and average GDP per capita

```
# load gapminder package
library(gapminder)
# load the dataset in object
gapminder = gapminder::gapminder
# display variables in the dataset
names(gapminder)
# show first 4 lines of the dataset
head(gapminder, n = 4)
```

```
## [1] "country" "continent" "year" "lifeExp" "pop" "gdpPercap"
## # A tibble: 4 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int> <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952  28.8  8425333  779.
## 2 Afghanistan Asia      1957  30.3  9240934  821.
## 3 Afghanistan Asia      1962  32.0 10267083  853.
## 4 Afghanistan Asia      1967  34.0 11537966  836.
```



In Practice: Data Wrangling

- There are always several ways to achieve a goal. (As in life 😊)
- Here we will only focus on the `dplyr` way:

```
# to install the dplyr package: install.packages("dplyr")
library(dplyr)
# compute the required statistics
gapminder_dplyr <- gapminder %>%
  group_by(continent, year) %>%
  summarise(
    count = n(),
    mean_lifeexp = mean(lifeExp, na.rm = TRUE),
    mean_gdppercap = mean(gdpPercap, na.rm = TRUE)
  )
```

```
# show first 5 lines of this dataframe
head(gapminder_dplyr, n = 5)
```


```
## # A tibble: 5 x 5
## # Groups:   continent [1]
##   continent year count mean_lifeexp mean_gdppercap
##   <fct>      <int> <int>      <dbl>      <dbl>
## 1 Africa    1952     52        39.1       1253.
## 2 Africa    1957     52        41.3       1385.
## 3 Africa    1962     52        43.3       1598.
## 4 Africa    1967     52        45.3       2050.
## 5 Africa    1972     52        47.5       2340.
```



Visualisation


- Now we could *look* at the result in `gapminder_dplyr`, or compute some statistics from it.
- Nothing beats a picture, though:

```
# to install the dplyr package: install.packages("dplyr")
library(ggplot2)
# create a plot
ggplot(data = gapminder_dplyr,
       mapping = aes(x = mean_lifeexp,
                     y = mean_gdppercap,
                     color = continent,
                     size = count)) +
  geom_point(alpha = 1/3) +
  labs(x = "Average life expectancy",
       y = "Average GDP per capita",
       color = "Continent",
       size = "Nb of countries")
```



Animated Plotting 🤞 1

Year: 1952



[1]: This animation is taken from [Ed Rubin](#).

R 101: Here Is Where You Start

Tool Time!

Getting R and Rstudio

- Download **R** from **CRAN** for your OS.
- Download **RStudio** from **here** for your OS.



Start your RStudio!

First Glossary of Terms

- R: a programming language.
- RStudio: an integrated development environment (IDE) to work with R.



Start your RStudio!

First Glossary of Terms

- **R**: a programming language.
- **RStudio**: an integrated development environment (IDE) to work with **R**.
- *command*: user input (text or numbers) that **R** *understands*.
- *script*: a list of commands collected in a text file, each separated by a new line, to be run one after the other.




Start your RStudio!

First Glossary of Terms

- **R**: a programming language.
- **RStudio**: an integrated development environment (IDE) to work with **R**.
- *command*: user input (text or numbers) that **R** *understands*.
- *script*: a list of commands collected in a text file, each separated by a new line, to be run one after the other.
- To run a script, you need to highlight the relevant code lines and hit **Ctrl+Enter** (Windows) or **Cmd+Enter** (Mac).



RStudio Layout



R as a Calculator

- You can use the `R` console like a calculator
- Just type an arithmetic operation after `>` and hit `Enter`!



R as a Calculator

- You can use the **R** console like a calculator
- Just type an arithmetic operation after **>** and hit **Enter**!

- Some basic arithmetic first:

```
4 + 1
```

```
## [1] 5
```

```
8 / 2
```

```
## [1] 4
```

- Great! What about this?

```
log(exp(1))
```

```
## [1] 1
```

```
# by the way: this is a comment! (R disregards it)
```



Task 1 (5 minutes)

- Create a new R script (File → New File → R Script). Save it somewhere as `lecture_intro.R`.
- Write in your script and run the following code: (Ctrl or Cmd + Enter)

```
4 * 8
```

- Write in your script and run the following code. What happens if you only run the first line of the code?

```
x = 5 # equivalently x <- 5  
x
```

- Congratulations, you have created your first R "object"! Everything is an object in R! Objects are assigned using `=` or `<-`.
- Lastly, find the cube of `x` and assign that value to `x_3`.



Where to get Help?

- R built-in help:

```
?log
?sin
?paste
?lm
help(lm)    # help() is equivalent
??plot     # get all help on keyword "plot"
help(ggplot,package="ggplot2") # show help from a certain package
```

- Help from Humans!
 - Google is your best friend!
 - stackoverflow.com [SO]
 - Your classroom channel on Slack
 - [rstudio forum](https://www.rstudio.com/forum)



HOW to get Help? Follow this for Slack questions!

1. Describe what you want to do.
2. Describe what you *expect* your code to do.
3. Describe what your code *does instead*.
 - Provide the entire error message.
4. Provide enough code to *reproduce* your error.
 - You can post code snippets on Slack and Stack Overflow



R Packages

- R users contribute add-on data and functions as *packages*
- Installing packages is easy!

```
install.packages("ggplot2")
```

- To *use* the contents of a package, we must load it from our library:

```
library(ggplot2)
```



ScPoEconometrics package

- The ScPoEconometrics team wrote an **R** package for you.
- It's hosted on **GitHub**
- You can install (and frequently update!) from here:

```
if (!require("devtools")) install.packages("devtools")  
library(devtools)  
install_github(repo = "ScPoEcon/ScPoEconometrics")
```



ScPoEconometrics package

- The ScPoEconometrics team wrote an R package for you.
- It's hosted on [GitHub](#)
- You can install (and frequently update!) from here:

```
if (!require("devtools")) install.packages("devtools")
library(devtools)
install_github(repo = "ScPoEcon/ScPoEconometrics")
```

- Did it work?

```
library(ScPoEconometrics)
packageVersion("ScPoEconometrics")
```

```
## [1] '0.2.6'
```



Vectors

- What is a **vector**?
- The **c** function creates vectors.

```
c(1, 3, 5, 7, 8, 9)
```

```
## [1] 1 3 5 7 8 9
```

- Coercion to unique types:

```
c(42, "Statistics", TRUE)
```

```
## [1] "42" "Statistics" "TRUE"
```

- Creating a *range*

```
c(y = 1:6)
```

```
## y1 y2 y3 y4 y5 y6
```

```
## 1 2 3 4 5 6
```



data.frame's

data.frames are like spreadsheets.

```
example_data = data.frame(x = c(1, 3, 5, 7),  
                           y = c(rep("Hello", 3), "Goodbye"),  
                           z = sample(c(TRUE, FALSE), size=4, replace=TRUE))
```

example_data

```
##   x      y      z  
## 1 1   Hello FALSE  
## 2 3   Hello FALSE  
## 3 5   Hello  TRUE  
## 4 7 Goodbye  TRUE
```

In practice, you will be importing files that contain the data into R rather than creating data.frames by hand.



Task 2 (10 minutes)

- Find out (using `help()` or google) how to import a .csv file.
- Import `gun_murders.csv`¹ in a new object `murders`. This file contains data on gun murders by US state in 2010. (Hint: objects are created using `=`)
- Ensure that `murders` is a data.frame by running:

```
# Check class  
class(murders)
```

- Find out what variables are contained in `murders` by running:

```
# Obtain variable names  
names(murders)
```

- View the contents of `murders` by clicking on `murders` in your workspace
 - What does the `total` variable correspond to?

[1]: This dataset is taken from the `dslabs` package.



data.frames

- Useful methods for a dataframe:

```
str(murders) # describes the data.frame
```

```
## 'data.frame':    51 obs. of  5 variables:
## $ state      : chr  "Alabama" "Alaska" "Arizona" "Arkansas" ...
## $ abb       : chr  "AL" "AK" "AZ" "AR" ...
## $ region    : Factor w/ 4 levels "Northeast","South",...: 2 4 4 2 4 4 1 2 2 2 ...
## $ population: num  4779736 710231 6392017 2915918 37253956 ...
## $ total     : num  135 19 232 93 1257 ...
```

```
names(murders) # column names
```

```
## [1] "state"      "abb"        "region"     "population" "total"
```

```
nrow(murders) # number of rows
```

```
## [1] 51
```

```
ncol(murders) # number of columns
```

```
## [1] 5
```



Data on Gun Murders in the US

- Let's **View** the data in a table format

```
View(murders) # Open a data.frame/vector in an Excel style table (up to 50 columns)
```

- Let's look at the first rows of `murders`.

```
head(murders, n = 3) # show first 3 rows
```

```
##      state abb region population total
## 1 Alabama  AL  South    4779736    135
## 2 Alaska   AK   West     710231     19
## 3 Arizona  AZ   West     6392017    232
```

- To access one of the variables **as a vector**, we use the `$` operator as in `murders$state`. We can check the type of `murders$state` with

```
class(murders$state) # type of the state variable in the murders data.frame
```

```
## [1] "character"
```

- Or we use the column name or index: `murders[, "state"]` or `murders[, 1]`



Subsetting data.frames

- Subsetting a data.frame: `murders[row condition, column number]` or `murders[row condition, "column name"]`

```
# Only keep states with over 500 gun murders and keep only the "state" and "total" variables  
murders[murders$total > 500, c("state", "total")]
```

```
##           state total  
## 5  California 1257  
## 10   Florida  669  
## 33   New York  517  
## 44    Texas   805
```

- There is also a special function for subsetting data:

```
subset(murders, subset = total > 500, select = c("state", "total"))
```

```
##           state total  
## 5  California 1257  
## 10   Florida  669  
## 33   New York  517  
## 44    Texas   805
```



Task 3 (10 minutes)

1. How many observations are there in `murders`?
2. How many variables? What are the data types of each variable?
3. Notice that the colon operator `a:b` is just short for *construct a sequence from `a` to `b`*. Create a new object `murders_2` containing the rows 10 to 25 of `murders`.
4. Create a new object `murders_3` which only contains the columns `state` and `total`. (Recall that `c` creates vectors.)
5. What is the average value of `total`?
6. What is the average value of `total` for state's in the "South", i.e. with `region == "South"`?
7. Create a `total_percap` variable:

```
murders$total_percap = (murders$total / murders$population) * 10000
```



Congratulations, you've created your first variable!

SEE YOU NEXT WEEK!

 michele.fioretti@sciencespo.fr

 Slides

 Book

 @ScPoEcon

 @ScPoEcon

