SciencesPo
DEPARTMENT OF ECONOMICS

Applied Data Analysis for Public Policy Studies

Regression Inference

Michele Fioretti
SciencesPo Paris
2020-08-25

Recap from last week

- ***Confidence interval***: a plausible range of value for the population parameter
- ***Hypothesis testing***: null hypothesis (H_0) vs alternative hypothesis (H_A), (observed) test statistic, null distribution
- ***p-value***: probability of observing a test statistic as or more extreme than the observed test statistic assuming the null hypothesis is true.

Recap from last week

- ***Confidence interval***: a plausible range of value for the population parameter
- ***Hypothesis testing***: null hypothesis (H_0) vs alternative hypothesis (H_A), (observed) test statistic, null distribution
- ***p-value***: probability of observing a test statistic as or more extreme than the observed test statistic assuming the null hypothesis is true.

Today: Statistical inference in the regression framework

- Fully understand a regression table
- Compare theory-based and simulation-based inference
- ***Classical Regression Model*** assumptions
- Empirical applications:
  - Class size and student performance
  - Returns to education by gender

Back to class size and student performance

- Let's go back the ***STAR*** experiment data, and focus on:
  - *small* and *regular* classes,
  - *Kindergarten* grade.

Back to class size and student performance

- Let's go back the ***STAR*** experiment data, and focus on:
  - *small* and *regular* classes,
  - *Kindergarten* grade.
- We consider the following regression model and estimate it by OLS:

$$\text{math\_score}_i = b_0 + b_1 \text{small}_i + e_i$$

Back to class size and student performance

- Let's go back the ***STAR*** experiment data, and focus on:
  - *small* and *regular* classes,
  - *Kindergarten* grade.
- We consider the following regression model and estimate it by OLS:

$$\text{math\_score}_i = b_0 + b_1 \text{small}_i + e_i$$

```
library(tidyverse)

star_df = read.csv("https://www.dropbox.com/s/bf1fog8y
star_df = star_df[complete.cases(star_df),]
star_df = star_df %>%
  filter(star %in% c("small","regular") &
           grade == "k") %>%
  mutate(small = (star == "small"))
```

Back to class size and student performance

- Let's go back the ***STAR*** experiment data, and focus on:
  - *small* and *regular* classes,
  - *Kindergarten* grade.
- We consider the following regression model and estimate it by OLS:

$$\text{math\_score}_i = b_0 + b_1 \text{small}_i + e_i$$

```
library(tidyverse)

star_df = read.csv("https://www.dropbox.com/s/bf1fog8y
star_df = star_df[complete.cases(star_df),]
star_df = star_df %>%
  filter(star %in% c("small","regular") &
           grade == "k") %>%
  mutate(small = (star == "small"))
```

```
lm(math ~ small, star_df)
```

```
##
## Call:
## lm(formula = math ~ small, data = star_df)
##
## Coefficients:
## (Intercept)    smallTRUE
##     484.446        8.895
```

Back to class size and student performance

- Let's go back the ***STAR*** experiment data, and focus on:
  - *small* and *regular* classes,
  - *Kindergarten* grade.
- We consider the following regression model and estimate it by OLS:

$$\text{math\_score}_i = b_0 + b_1 \text{small}_i + e_i$$

```
library(tidyverse)

star_df = read.csv("https://www.dropbox.com/s/bf1fog8y
star_df = star_df[complete.cases(star_df),]
star_df = star_df %>%
  filter(star %in% c("small","regular") &
           grade == "k") %>%
  mutate(small = (star == "small"))
```

```
lm(math ~ small, star_df)
```

```
##
## Call:
## lm(formula = math ~ small, data = star_df)
##
## Coefficients:
## (Intercept)    smallTRUE
##     484.446        8.895
```

- What if we drew another random sample of schools from Tennessee and redid the experiment, would we find a different value for b_1?
- We know the answer is *yes*, but how different is this estimate likely to be?

Regression Inference: b_k vs β_k

- b_0, b_1 are ***point estimates*** computed from our sample.
  - Just like the sample proportion $\hat{p}$ from our pasta example!

Regression Inference: b_k vs β_k

- b_0, b_1 are ***point estimates*** computed from our sample.
  - Just like the sample proportion $\hat{p}$ from our pasta example!

- In fact, our model's prediction...

$$\hat{y} = b_0 + b_1 x_1$$

Regression Inference: b_k vs β_k

- b_0, b_1 are ***point estimates*** computed from our sample.
  - Just like the sample proportion $\hat{p}$ from our pasta example!

- In fact, our model's prediction...

$$\hat{y} = b_0 + b_1 x_1$$

... is an **estimate** about an unknown, **true population line**

$$y = \beta_0 + \beta_1 x_1$$

where β_0, β_1 are the ***population parameters*** of interest.

Regression Inference: b_k vs β_k

- b_0, b_1 are ***point estimates*** computed from our sample.
  - Just like the sample proportion $\hat{p}$ from our pasta example!

- In fact, our model's prediction...

$$\hat{y} = b_0 + b_1 x_1$$

... is an **estimate** about an unknown, **true population line**

$$y = \beta_0 + \beta_1 x_1$$

where β_0, β_1 are the ***population parameters*** of interest.

- You will often find $\hat{\beta}_k$ rather than b_k, both refer to sample estimate of β_k.

Regression Inference: b_k vs β_k

- b_0, b_1 are ***point estimates*** computed from our sample.
  - Just like the sample proportion $\hat{p}$ from our pasta example!

- In fact, our model's prediction...

$$\hat{y} = b_0 + b_1 x_1$$

... is an **estimate** about an unknown, **true population line**

$$y = \beta_0 + \beta_1 x_1$$

where β_0, β_1 are the ***population parameters*** of interest.

- You will often find $\hat{\beta}_k$ rather than b_k, both refer to sample estimate of β_k.
- Let's bring what we know about ***confidence intervals***, ***hypothesis testing*** and ***standard errors*** to bear on those $\hat{\beta}_k$!

Understanding Regression Tables

Here is our `tidy` regression:

```
library(broom)
tidy(lm(math ~ small, star_df))
```

```
## # A tibble: 2 x 5
##   term        estimate std.error statistic     p.value
##   <chr>          <dbl>     <dbl>     <dbl>       <dbl>
## 1 (Intercept)   484.        1.15    421.    0
## 2 smallTRUE       8.90      1.68      5.30 0.000000123
```

- There are 3 new columns here: `std.error`, `statistic`, `p.value`.

Understanding Regression Tables

Here is our `tidy` regression:

```
library(broom)
tidy(lm(math ~ small, star_df))
```

```
## # A tibble: 2 x 5
##   term        estimate std.error statistic     p.value
##   <chr>          <dbl>     <dbl>     <dbl>       <dbl>
## 1 (Intercept)   484.        1.15    421.    0
## 2 smallTRUE       8.90      1.68      5.30 0.000000123
```

- There are 3 new columns here: `std.error`, `statistic`, `p.value`.

Entry	Meaning
`std. error`	Standard error of b_k
`statistic`	Observed test statistic associated to $H_0 : \beta_k = 0, H_A : \beta_k \neq 0$
`p.value`	p-value associated to $H_0 : \beta_k = 0, H_A : \beta_k \neq 0$

Understanding Regression Tables

Here is our `tidy` regression:

```
library(broom)
tidy(lm(math ~ small, star_df))
```

```
## # A tibble: 2 x 5
##   term        estimate std.error statistic     p.value
##   <chr>          <dbl>     <dbl>     <dbl>       <dbl>
## 1 (Intercept)   484.        1.15    421.    0
## 2 smallTRUE       8.90      1.68      5.30 0.000000123
```

- There are 3 new columns here: `std.error`, `statistic`, `p.value`.

Entry	Meaning
`std. error`	Standard error of b_k
`statistic`	Observed test statistic associated to $H_0 : \beta_k = 0, H_A : \beta_k \neq 0$
`p.value`	p-value associated to $H_0 : \beta_k = 0, H_A : \beta_k \neq 0$

- Let's focus on the `small` coefficient and make sense of each entry.

Standard Error of b_k

> ***Standard Error of b_k:*** Standard deviation of the sampling distribution of b_k.

Standard Error of b_k

> ***Standard Error of b_k:*** Standard deviation of the sampling distribution of b_k.

Let's imagine we could redo the experiment 1000 times on 1000 different samples:

- We'd run 1000 regression and would get 1000 estimates of β_k, b_k.

Standard Error of b_k

> ***Standard Error of b_k:*** Standard deviation of the sampling distribution of b_k.

Let's imagine we could redo the experiment 1000 times on 1000 different samples:

- We'd run 1000 regression and would get 1000 estimates of β_k, b_k.
- The standard error of b_k quantifies how much variation in b_k one would expect across (*an infinity of*) samples.

Standard Error of b_{small}

- From the table, we get $\hat{\text{SE}}(b_{\text{small}}) = 1.68$
    - Notice that we write $\hat{\text{SE}}$ and not SE because 1.68 is an estimate of the real standard error of b_{small} we get from our sample.

Standard Error of b_{small}

- From the table, we get $\hat{\text{SE}}(b_{\text{small}}) = 1.68$
    - Notice that we write $\hat{\text{SE}}$ and not SE because 1.68 is an estimate of the real standard error of b_{small} we get from our sample.
- Let's simulate the sampling distribution of b_{small} to see where it comes from.

Task 1 (10 min)

As we did for the sampling distribution of the proportion of *green pasta*, we want to generate the bootstrap distribution of b_{small}.

1. Copy the loading and cleaning code from slide 3 and run it.

2. Generate the bootstrap distribution of b_{small} based on 1000 samples drawn from `star_df`. *Hint*: use the appropriate functions and arguments from the `infer` package so use the help pages.

3. Plot this simulated sampling distribution and compute mean and the standard error of b_{small}.

Bootstrap Distribution

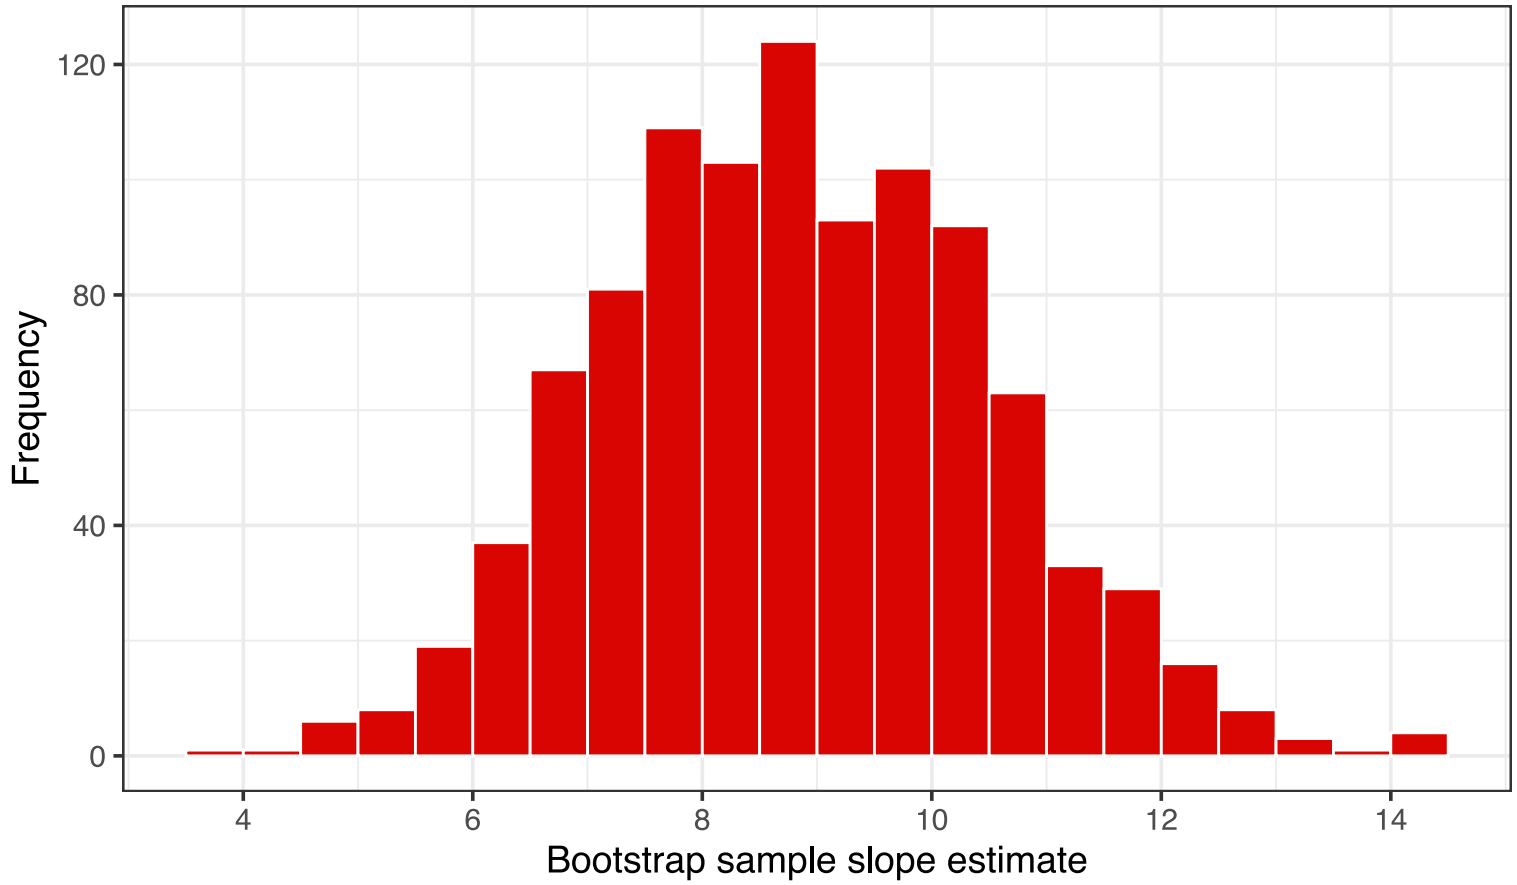

Bootstrap Distribution

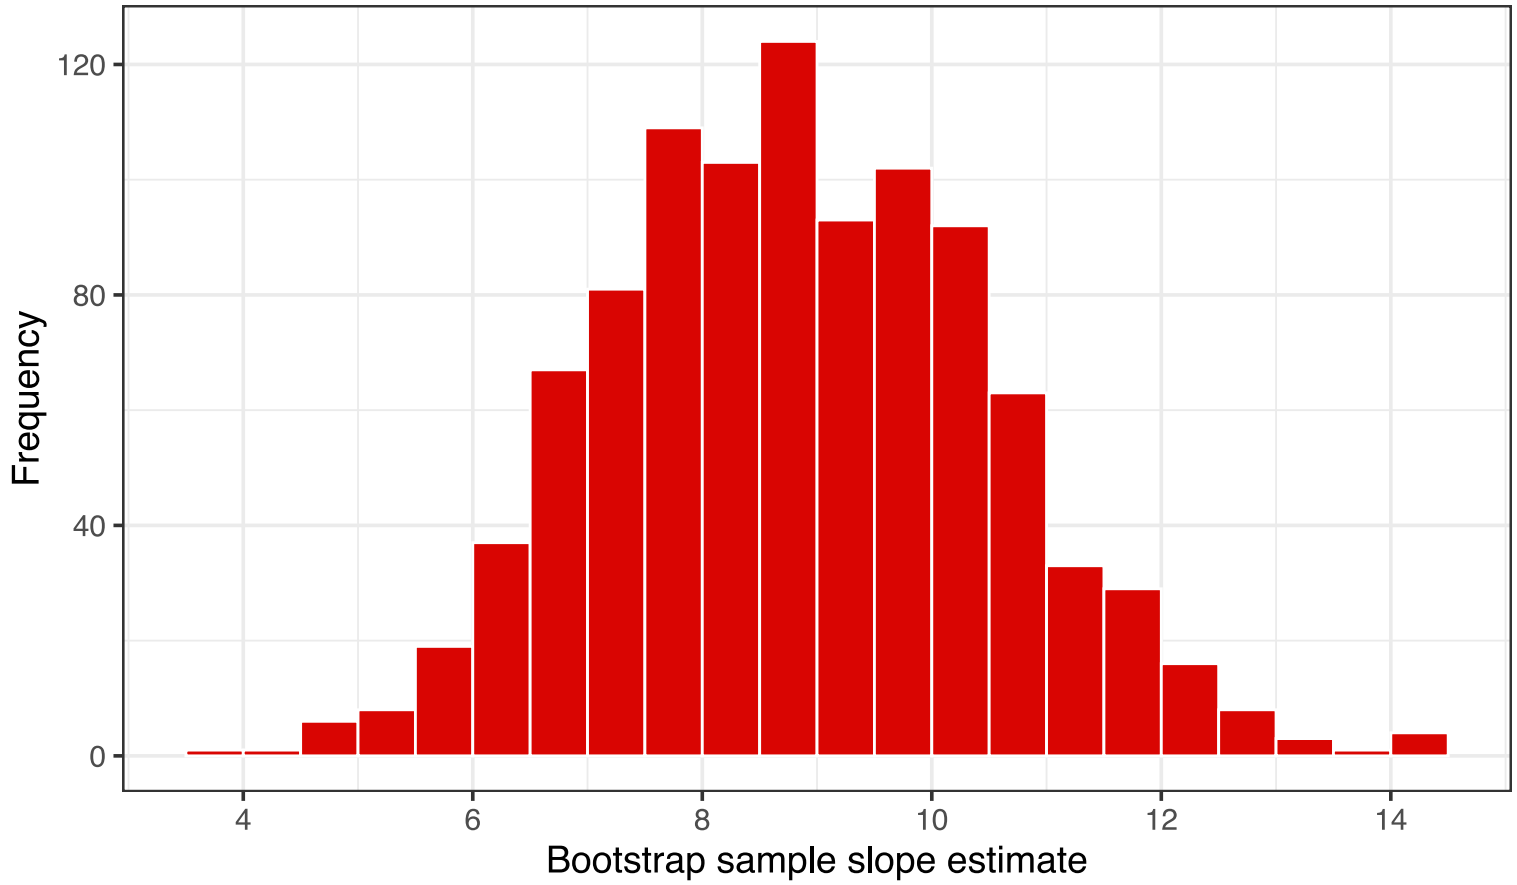

standard error: 1.67 $\rightarrow$ very close to the one in the table (1.68)!

Testing $\beta_k = 0$ vs $\beta_k \neq 0$

By default, the regression output provides the results associated with the following hypothesis test:

$$H_0 : \beta_k = 0$$
$$H_A : \beta_k \neq 0$$

Testing $\beta_k = 0$ vs $\beta_k \neq 0$

By default, the regression output provides the results associated with the following hypothesis test:

$$H_0 : \beta_k = 0$$
$$H_A : \beta_k \neq 0$$

- It allows to statistically test if there is a true relationship between the outcome and our regressor.

Testing $\beta_k = 0$ vs $\beta_k \neq 0$

By default, the regression output provides the results associated with the following hypothesis test:

$$H_0 : \beta_k = 0$$
$$H_A : \beta_k \neq 0$$

- It allows to statistically test if there is a true relationship between the outcome and our regressor.
- If H_0 is true, there is **no** relationship between the outcome and our regressor.
  - In that case observing $b_1 \neq 0$ was just chance.

Testing $\beta_k = 0$ vs $\beta_k \neq 0$

By default, the regression output provides the results associated with the following hypothesis test:

$$H_0 : \beta_k = 0$$
$$H_A : \beta_k \neq 0$$

- It allows to statistically test if there is a true relationship between the outcome and our regressor.
- If H_0 is true, there is **no** relationship between the outcome and our regressor.
  - In that case observing $b_1 \neq 0$ was just chance.
- If H_0 is false, then there **is** a true relationship.

Testing $\beta_k = 0$ vs $\beta_k \neq 0$

By default, the regression output provides the results associated with the following hypothesis test:

$$H_0 : \beta_k = 0$$
$$H_A : \beta_k \neq 0$$

- It allows to statistically test if there is a true relationship between the outcome and our regressor.
- If H_0 is true, there is **no** relationship between the outcome and our regressor.
  - In that case observing $b_1 \neq 0$ was just chance.
- If H_0 is false, then there **is** a true relationship.
- ***Important:*** This is a ***two-sided*** test!

Test statistic and p-value

- As we saw in the previous lecture, to conduct such a test we need to:

Test statistic and p-value

- As we saw in the previous lecture, to conduct such a test we need to:
  - Derive the sampling distribution of our **test statistic** (`statistic`) assuming H_0 is true, i.e. the *null distribution*.

Test statistic and p-value

- As we saw in the previous lecture, to conduct such a test we need to:
  - Derive the sampling distribution of our **test statistic** (`statistic`) assuming H_0 is true, i.e. the *null distribution*.
  - Quantify how extreme the **observed test statistic** is in this hypothetic world.

Test statistic and p-value

- As we saw in the previous lecture, to conduct such a test we need to:
  - Derive the sampling distribution of our **test statistic** (`statistic`) assuming H_0 is true, i.e. the *null distribution*.
  - Quantify how extreme the **observed test statistic** is in this hypothetic world.
- Our *observed test statistic* (`statistic`) equals $\frac{b}{\hat{SE}(b)}$.
  - Why not just b? We'll come back and explain this formula later.

Test statistic and p-value

- As we saw in the previous lecture, to conduct such a test we need to:
  - Derive the sampling distribution of our **test statistic** (`statistic`) assuming H_0 is true, i.e. the *null distribution*.
  - Quantify how extreme the **observed test statistic** is in this hypothetic world.
- Our *observed test statistic* (`statistic`) equals $\frac{b}{\hat{SE}(b)}$.
  - Why not just b? We'll come back and explain this formula later.

```
observed_stat = lm(math ~ small, star_df)$coefficients
round(observed_stat,2)
```

```
## smallTRUE
##      5.33
```

Test statistic and p-value

- As we saw in the previous lecture, to conduct such a test we need to:
  - Derive the sampling distribution of our **test statistic** (`statistic`) assuming H_0 is true, i.e. the *null distribution*.
  - Quantify how extreme the **observed test statistic** is in this hypothetic world.
- Our *observed test statistic* (`statistic`) equals $\frac{b}{\hat{SE}(b)}$.
  - Why not just b? We'll come back and explain this formula later.

```
observed_stat = lm(math ~ small, star_df)$coefficients
round(observed_stat,2)
```

```
## smallTRUE
##      5.33
```

- Quite close to the observed test statistic we got in the table: `statistic` = 5.3.

Test statistic and p-value

- As we saw in the previous lecture, to conduct such a test we need to:
  - Derive the sampling distribution of our **test statistic** (`statistic`) assuming H_0 is true, i.e. the *null distribution*.
  - Quantify how extreme the **observed test statistic** is in this hypothetic world.
- Our *observed test statistic* (`statistic`) equals $\frac{b}{\hat{SE}(b)}$.
  - Why not just b? We'll come back and explain this formula later.

```
observed_stat = lm(math ~ small, star_df)$coefficients
round(observed_stat,2)
```

```
## smallTRUE
##      5.33
```

- Quite close to the observed test statistic we got in the table: `statistic` = 5.3.

- The **p-value** measures the area outside of $\pm$ *observed test statistic* under the *null distribution*.

Test statistic and p-value

- As we saw in the previous lecture, to conduct such a test we need to:
  - Derive the sampling distribution of our **test statistic** (`statistic`) assuming H_0 is true, i.e. the *null distribution*.
  - Quantify how extreme the **observed test statistic** is in this hypothetic world.
- Our *observed test statistic* (`statistic`) equals $\frac{b}{\hat{SE}(b)}$.
  - Why not just b? We'll come back and explain this formula later.

```
observed_stat = lm(math ~ small, star_df)$coefficients
round(observed_stat,2)
```

```
## smallTRUE
##      5.33
```

- Quite close to the observed test statistic we got in the table: `statistic` = 5.3.

- The **p-value** measures the area outside of $\pm$ *observed test statistic* under the *null distribution*.
- Finally, we check if we can reject H_0 at the usual **significance levels**: α = 0.1, 0.05, 0.01.

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

- We will approximate the null distribution of $\frac{b_{\text{small}}}{\hat{SE}(b_{\text{small}})}$ through a simulation exercise.

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

- We will approximate the null distribution of $\frac{b_{\text{small}}}{\hat{SE}(b_{\text{small}})}$ through a simulation exercise.
- If there is no relationship between math score and class size, i.e. H_0 is true, then *reshuffling* / *permuting* the values of `small` across students should play no role.

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

- We will approximate the null distribution of $\frac{b_{\text{small}}}{\hat{SE}(b_{\text{small}})}$ through a simulation exercise.
- If there is no relationship between math score and class size, i.e. H_0 is true, then *reshuffling* / *permuting* the values of `small` across students should play no role.
- Let's generate 1000 permuted samples and compute b_{small} for each.

```
set.seed(123)
null_distribution <- star_df %>%
  specify(formula = math ~ small) %>%
  hypothesize(null = "independence") %>%
  generate(reps = 1000, type = "permute") %>%
  calculate(stat = "slope", order = c("TRUE","FALSE"))
```

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

- We will approximate the null distribution of $\frac{b_{\text{small}}}{\hat{SE}(b_{\text{small}})}$ through a simulation exercise.
- If there is no relationship between math score and class size, i.e. H_0 is true, then *reshuffling* / *permuting* the values of `small` across students should play no role.
- Let's generate 1000 permuted samples and compute b_{small} for each.

```
set.seed(123)
null_distribution <- star_df %>%
  specify(formula = math ~ small) %>%
  hypothesize(null = "independence") %>%
  generate(reps = 1000, type = "permute") %>%
  calculate(stat = "slope", order = c("TRUE","FALSE"))
```

- We can compute the distribution of our test statistic $\frac{b_{\text{small}}}{\hat{SE}(b_{\text{small}})}$ under the null:

```
null_distribution <- null_distribution %>%
  mutate(test_stat = stat/sd(bootstrap_distrib$stat))
```

- Remember we got $\hat{SE}(b_{\text{small}})$ = 1.67 from our bootstrap distribution.

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

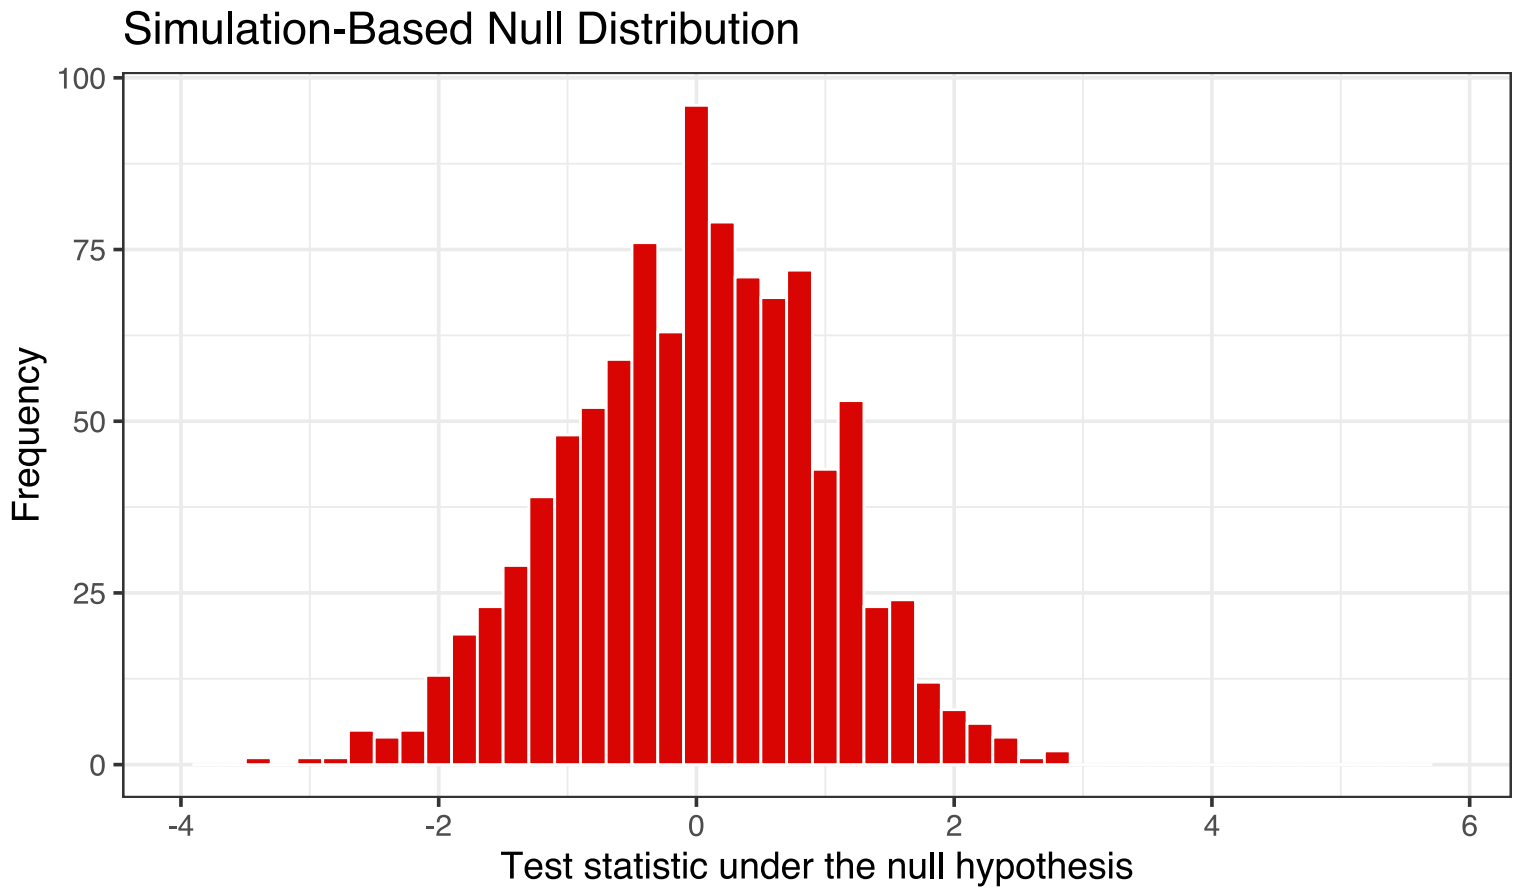

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

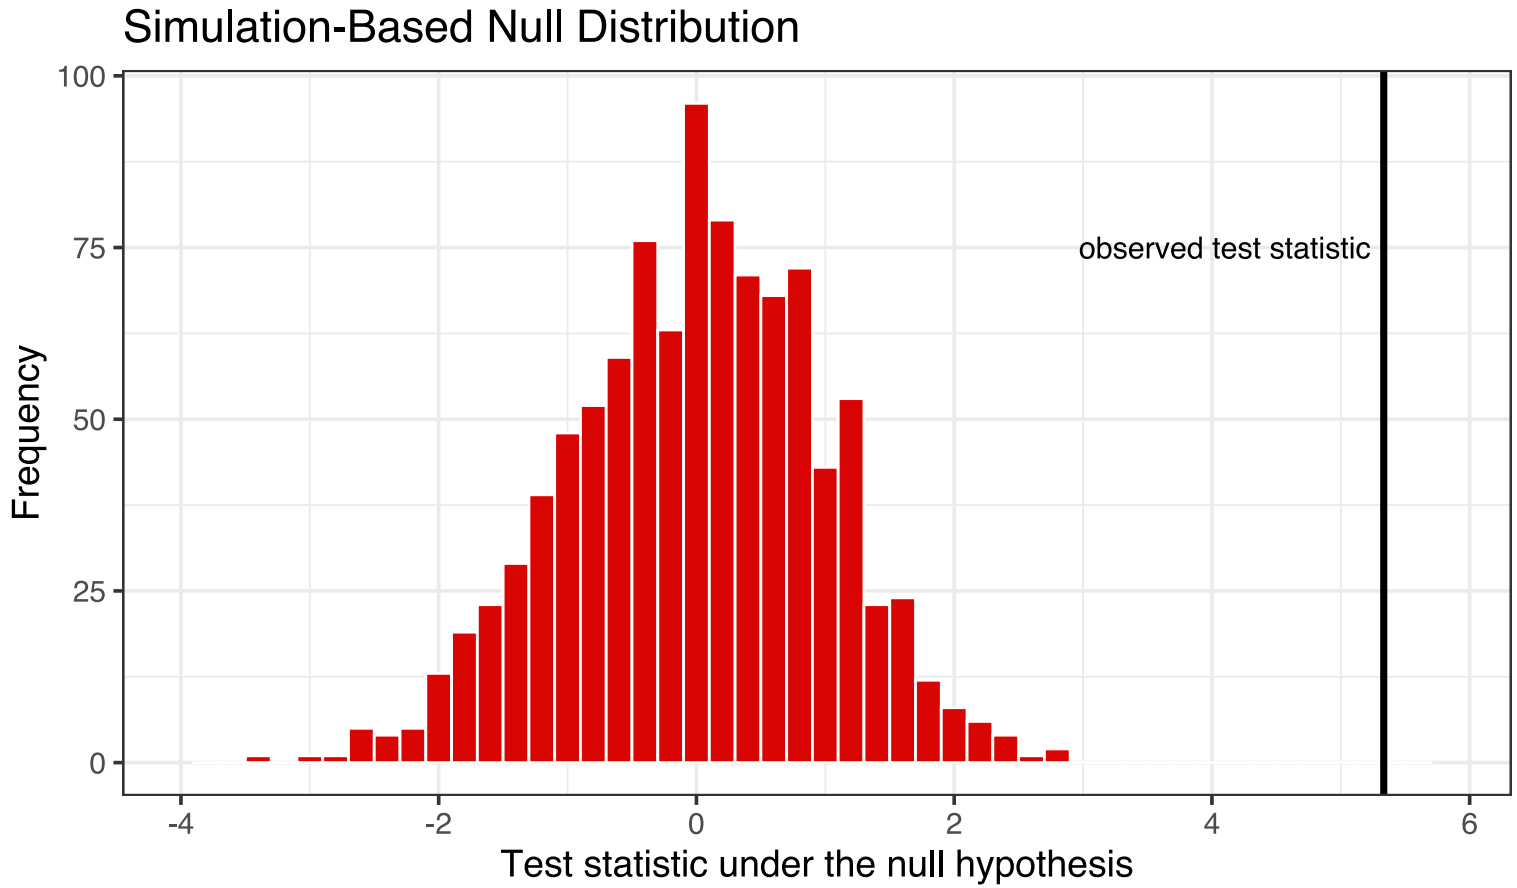

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

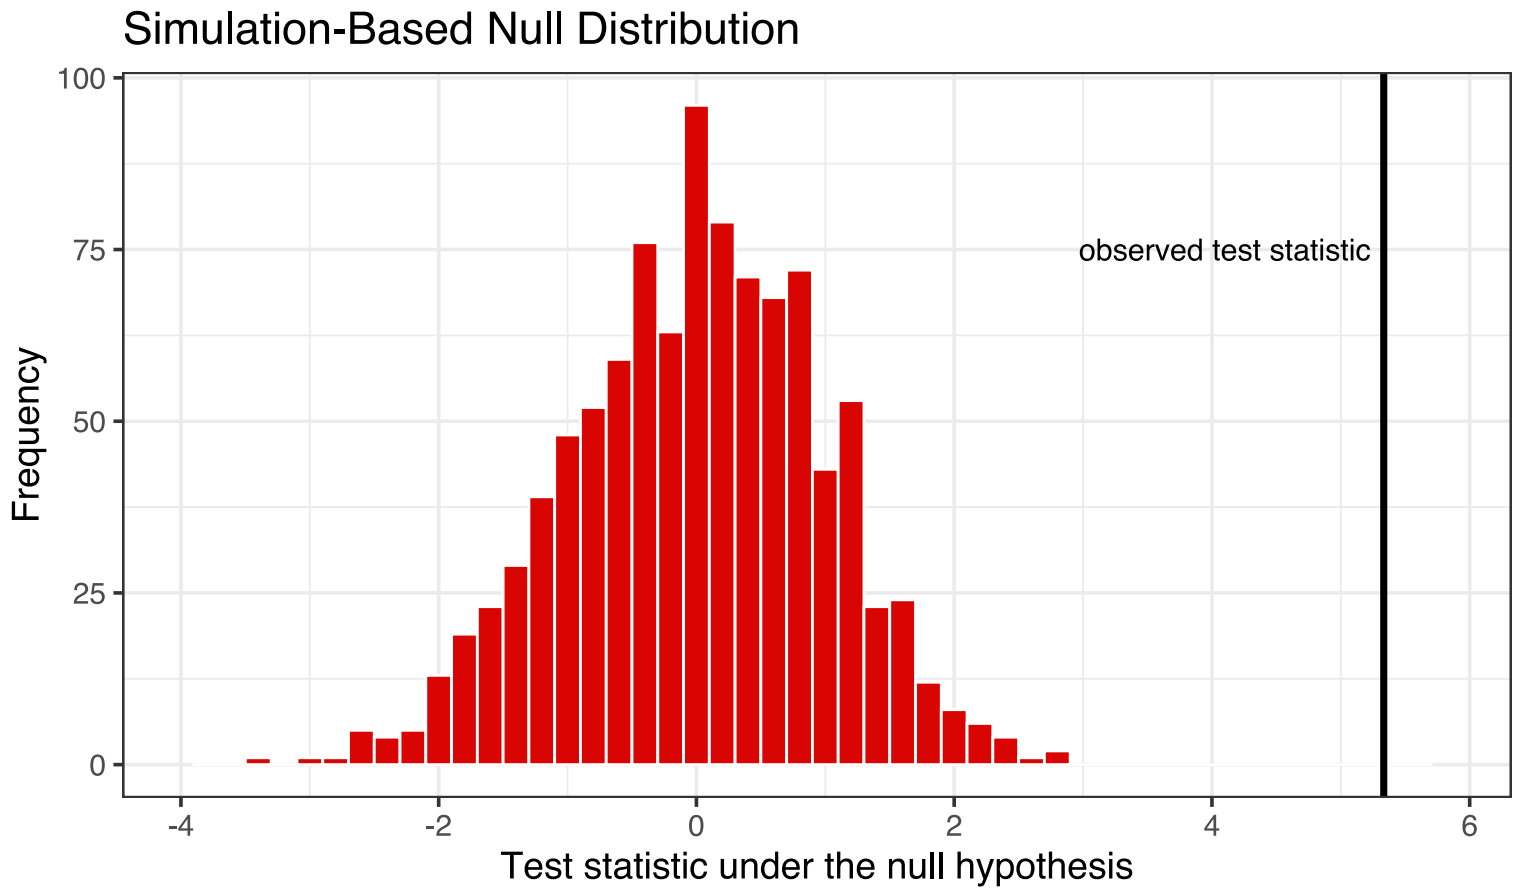

Very unlikely to obtain b_{small} = 8.9 when H_0 is true.

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

- To decide if we reject H_0, recall we are considering a **two-sided test** here: *more extreme* means inferior to -5.333 **or** superior to 5.333.

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

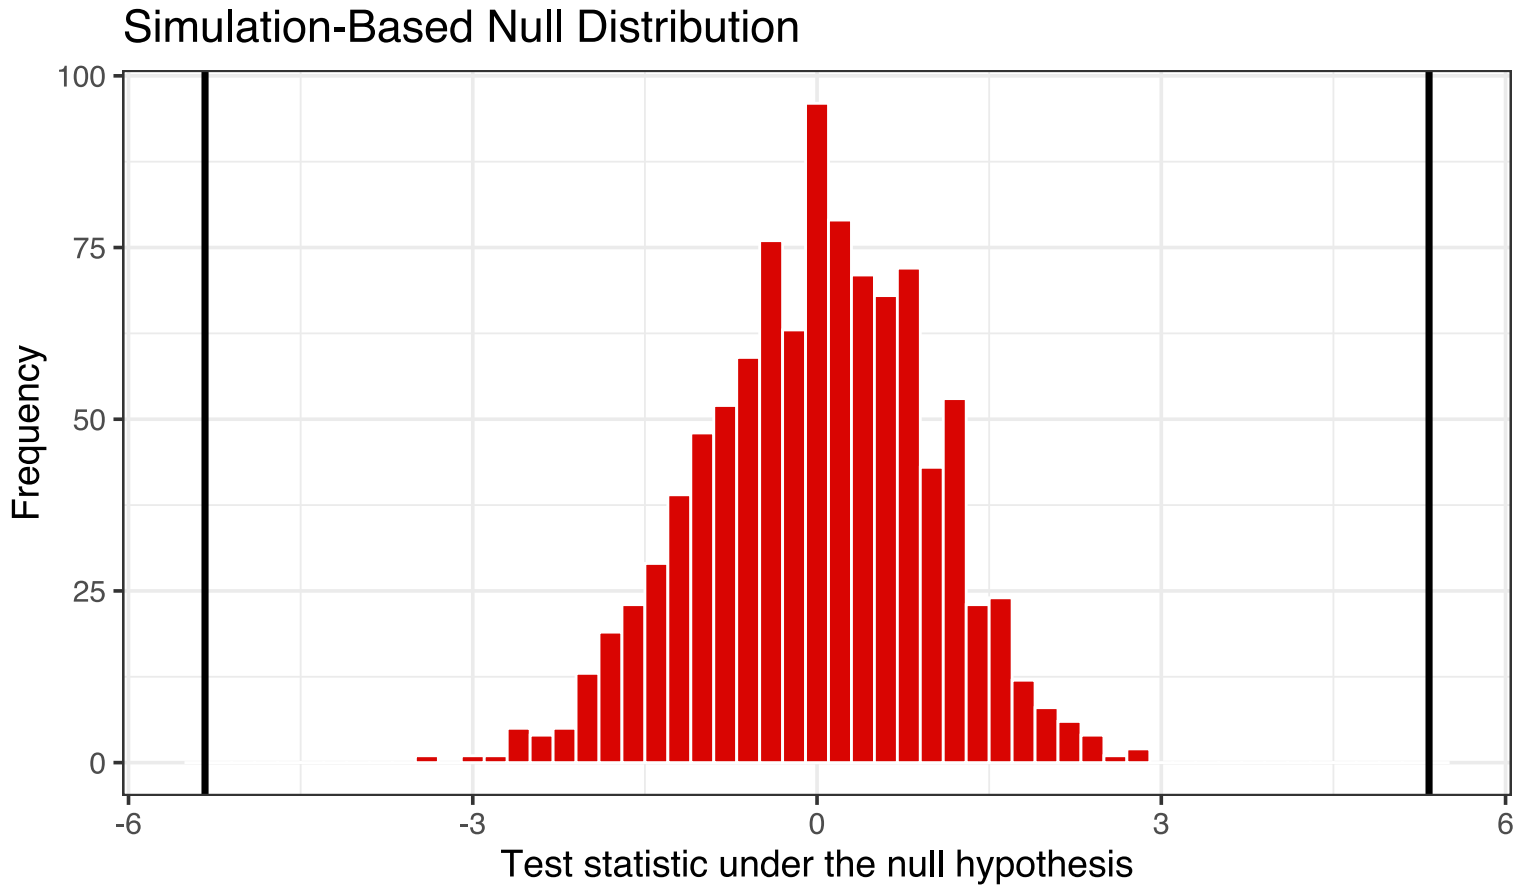

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

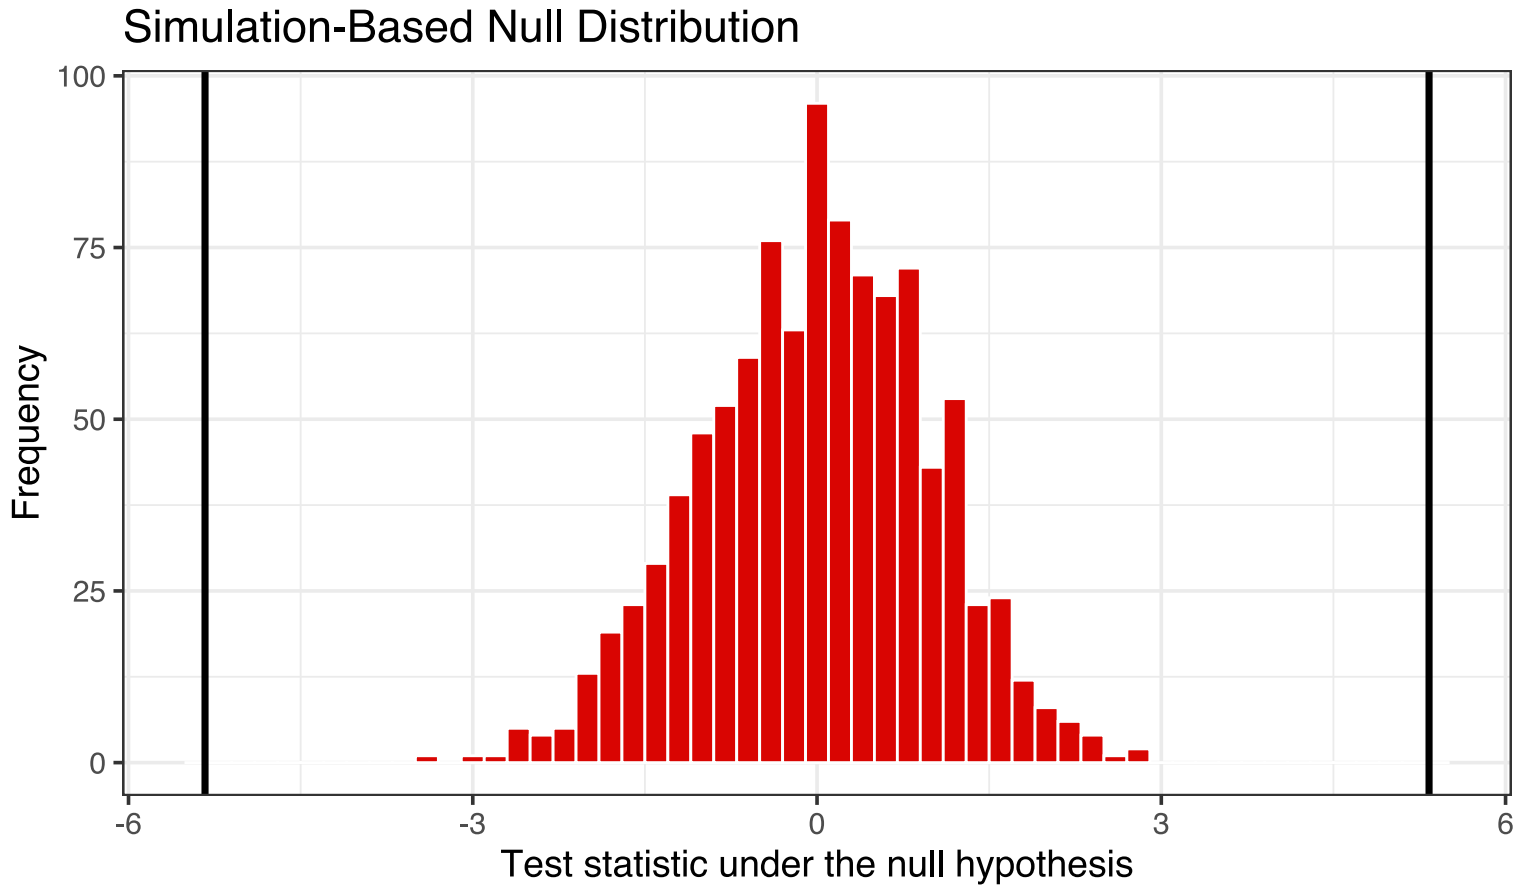

What does the p-value correspond to?

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

- To decide if we reject H_0, recall we are considering a **two-sided test** here: *more extreme* means inferior to -5.33 **or** superior to 5.33.
- Computing the *p-value* we get:

```
p_value = mean(abs(null_distribution$test_stat) >= observed_stat)
p_value
```

```
## [1] 0
```

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

- To decide if we reject H_0, recall we are considering a **two-sided test** here: *more extreme* means inferior to -5.33 **or** superior to 5.33.
- Computing the *p-value* we get:

```
p_value = mean(abs(null_distribution$test_stat) >= observed_stat)
p_value
```

```
## [1] 0
```

- This is the same value as in the regression table.

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

- To decide if we reject H_0, recall we are considering a **two-sided test** here: *more extreme* means inferior to -5.33 **or** superior to 5.33.
- Computing the *p-value* we get:

```
p_value = mean(abs(null_distribution$test_stat) >= observed_stat)
p_value
```

```
## [1] 0
```

- This is the same value as in the regression table.
- ***Question:*** Can we reject the null hypothesis at the 5% level?

Testing $\beta_{\text{small}} = 0$ vs $\beta_{\text{small}} \neq 0$

- To decide if we reject H_0, recall we are considering a **two-sided test** here: *more extreme* means inferior to -5.33 **or** superior to 5.33.
- Computing the *p-value* we get:

```
p_value = mean(abs(null_distribution$test_stat) >= observed_stat)
p_value
```

```
## [1] 0
```

- This is the same value as in the regression table.
- ***Answer:***
  - Since the *p-value* is equal to 0 it means that we would reject H_0 at any significance level: the p-value would always be inferior to α.
  - In other words, we can say that b_{small} is **statistically different from 0** at any significance level.
  - We also say that b_{small} is *statistically significant* (at any significance level).

Regression Inference: Theory

Regression Inference: Theory

- Up to now we presented simulation-based inference.

Regression Inference: Theory

- Up to now we presented simulation-based inference.
- The values reported by statistical packages in `R` are instead obtained from theory.

Regression Inference: Theory

- Up to now we presented simulation-based inference.
- The values reported by statistical packages in `R` are instead obtained from theory.
- Theoretical inference is based on **large sample approximations**.
    - One can show that sampling distributions *converge* to suitable distributions.

Regression Inference: Theory

- Up to now we presented simulation-based inference.
- The values reported by statistical packages in `R` are instead obtained from theory.
- Theoretical inference is based on **large sample approximations**.
  - One can show that sampling distributions *converge* to suitable distributions.
- Let's briefly look into the theory-based approach.

Regression Inference: Theory

- Theory-based approach uses one fundamental result: the sampling distribution of the sample statistic $\frac{b-\beta}{\hat{\text{SE}}(b)}$ *converges* to a ***standard normal distribution*** as the sample size gets larger and larger.
    - $\hat{\text{SE}}(b)$ is the sample estimate of the standard deviation of b.
    - It is also obtained through a theoretical formula (which you can find in the book!) but we'll leave it aside.

Regression Inference: Theory

- Theory-based approach uses one fundamental result: the sampling distribution of the sample statistic $\frac{b-\beta}{\hat{\text{SE}}(b)}$ *converges* to a ***standard normal distribution*** as the sample size gets larger and larger.

  - $\hat{\text{SE}}(b)$ is the sample estimate of the standard deviation of b.
  - It is also obtained through a theoretical formula (which you can find in the book!) but we'll leave it aside.
- A ***standard normal distribution*** is a *normal distribution* with *mean* 0 and *standard deviation* 1.

Regression Inference: Theory

- Theory-based approach uses one fundamental result: the sampling distribution of the sample statistic $\frac{b-\beta}{\hat{\text{SE}}(b)}$ *converges* to a ***standard normal distribution*** as the sample size gets larger and larger.
  - $\hat{\text{SE}}(b)$ is the sample estimate of the standard deviation of b.
  - It is also obtained through a theoretical formula (which you can find in the book!) but we'll leave it aside.
- A ***standard normal distribution*** is a *normal distribution* with *mean* 0 and *standard deviation* 1.
- We don't need to simulate any sampling distribution here, we derive it from theory and use it to construct confidence intervals or to conduct hypothesis tests.

Regression Inference: Theory

- Theory-based approach uses one fundamental result: the sampling distribution of the sample statistic $\frac{b-\beta}{\hat{\text{SE}}(b)}$ *converges* to a ***standard normal distribution*** as the sample size gets larger and larger.
  - $\hat{\text{SE}}(b)$ is the sample estimate of the standard deviation of b.
  - It is also obtained through a theoretical formula (which you can find in the book!) but we'll leave it aside.
- A ***standard normal distribution*** is a *normal distribution* with *mean* 0 and *standard deviation* 1.
- We don't need to simulate any sampling distribution here, we derive it from theory and use it to construct confidence intervals or to conduct hypothesis tests.
- Note that if $\frac{b-\beta}{\hat{\text{SE}}(b)}$ *converges* to a ***standard normal distribution***, then b converges to a ***normal distribution*** with mean β and standard deviation $\hat{\text{SE}}(b)$.

Standard Normal Distribution: A Refresher

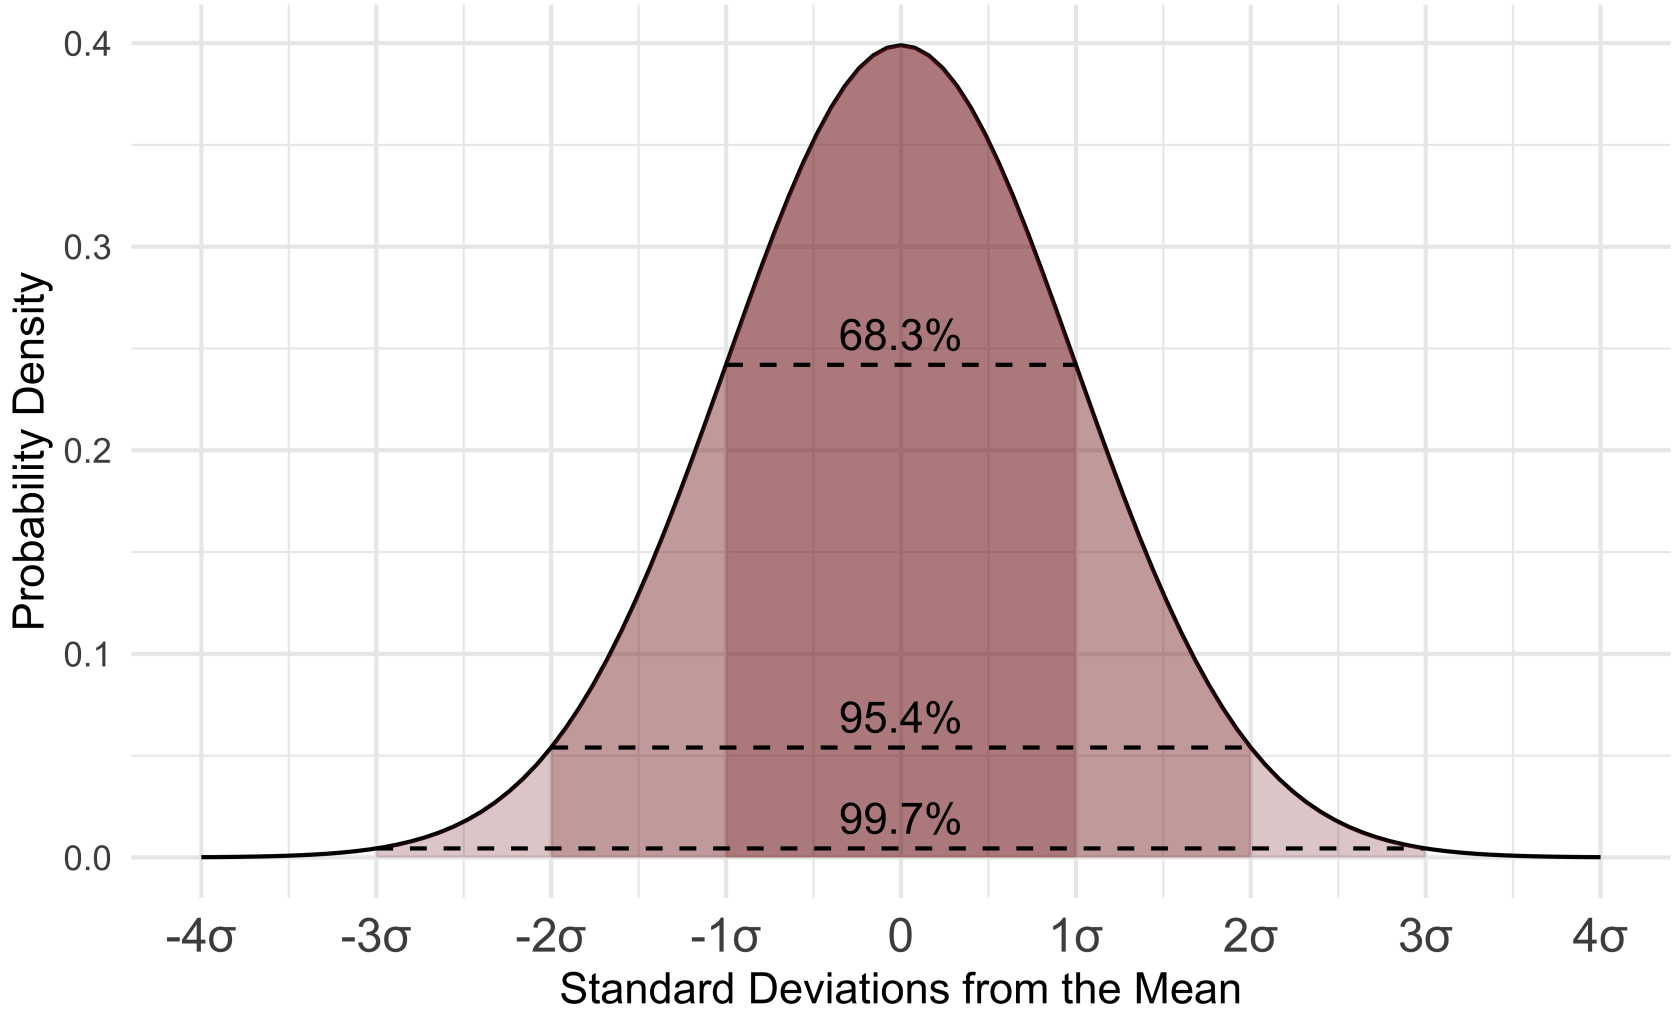

Theory-Based Inference: Confidence Interval

- Let's take the example of a 95% confidence interval.

Theory-Based Inference: Confidence Interval

- Let's take the example of a 95% confidence interval.
- Since the sampling distribution of b is assumed to be normally shaped, we can use the ***95% rule of thumb*** about normal distributions.

Theory-Based Inference: Confidence Interval

- Let's take the example of a 95% confidence interval.
- Since the sampling distribution of b is assumed to be normally shaped, we can use the ***95% rule of thumb*** about normal distributions.
- We know indeed that 95% of the values of a normal distribution lie within approximately 2 standard deviations of the mean (exactly 1.96).

Theory-Based Inference: Confidence Interval

- Let's take the example of a 95% confidence interval.
- Since the sampling distribution of b is assumed to be normally shaped, we can use the ***95% rule of thumb*** about normal distributions.
- We know indeed that 95% of the values of a normal distribution lie within approximately 2 standard deviations of the mean (exactly 1.96).
- So, we can compute a 95% CI for β as: $\mathrm{CI}_{95\%} = [b \pm 1.96 * \hat{\mathrm{SE}}(b)]$

Theory-Based Inference: Confidence Interval

- Let's take the example of a 95% confidence interval.
- Since the sampling distribution of b is assumed to be normally shaped, we can use the ***95% rule of thumb*** about normal distributions.
- We know indeed that 95% of the values of a normal distribution lie within approximately 2 standard deviations of the mean (exactly 1.96).
- So, we can compute a 95% CI for β as: $\text{CI}_{95\%} = [b \pm 1.96 * \hat{\text{SE}}(b)]$

```
tidy(lm(math ~ small, star_df),
     conf.int = TRUE, conf.level = 0.95) %>% # To dis
  filter(term == "smallTRUE") %>%
  select(term, conf.low, conf.high)
```

```
## # A tibble: 1 x 3
##   term      conf.low conf.high
##   <chr>        <dbl>     <dbl>
## 1 smallTRUE     5.60      12.2
```

Theory-Based Inference: Confidence Interval

- Let's take the example of a 95% confidence interval.
- Since the sampling distribution of b is assumed to be normally shaped, we can use the ***95% rule of thumb*** about normal distributions.
- We know indeed that 95% of the values of a normal distribution lie within approximately 2 standard deviations of the mean (exactly 1.96).
- So, we can compute a 95% CI for β as: $\text{CI}_{95\%} = [b \pm 1.96 * \hat{\text{SE}}(b)]$

```
tidy(lm(math ~ small, star_df),
     conf.int = TRUE, conf.level = 0.95) %>% # To dis
  filter(term == "smallTRUE") %>%
  select(term, conf.low, conf.high)
```

```
## # A tibble: 1 x 3
##   term      conf.low conf.high
##   <chr>        <dbl>     <dbl>
## 1 smallTRUE     5.60      12.2
```

```
bootstrap_distrib %>%
  summarise(
    lower_bound = 8.895 - 1.96*sd(stat),
    upper_bound = 8.895 + 1.96*sd(stat))
```

```
## # A tibble: 1 x 2
##   lower_bound upper_bound
##         <dbl>       <dbl>
## 1        5.63        12.2
```

Theory-Based Inference: Confidence Interval

- Let's take the example of a 95% confidence interval.
- Since the sampling distribution of b is assumed to be normally shaped, we can use the ***95% rule of thumb*** about normal distributions.
- We know indeed that 95% of the values of a normal distribution lie within approximately 2 standard deviations of the mean (exactly 1.96).
- So, we can compute a 95% CI for β as: $\text{CI}_{95\%} = [b \pm 1.96 * \hat{\text{SE}}(b)]$

```
tidy(lm(math ~ small, star_df),
     conf.int = TRUE, conf.level = 0.95) %>% # To dis
  filter(term == "smallTRUE") %>%
  select(term, conf.low, conf.high)
```

```
## # A tibble: 1 x 3
##   term      conf.low conf.high
##   <chr>        <dbl>     <dbl>
## 1 smallTRUE     5.60      12.2
```

```
bootstrap_distrib %>%
  summarise(
    lower_bound = 8.895 - 1.96*sd(stat),
    upper_bound = 8.895 + 1.96*sd(stat))
```

```
## # A tibble: 1 x 2
##   lower_bound upper_bound
##         <dbl>       <dbl>
## 1        5.63        12.2
```

- This can easily be generalized to any confidence level by taking the appropriate quantile of the normal distribution.

Task 2 (5 min)

1. Using the bootstrap distributed you generated in Task 1, compute the 95% confidence interval using the *percentile method*.

2. How similar is it to the confidence intervals obtained in the previous slide?

Theory-Based Inference: Hypothesis Testing

- As we already mentioned, the default test that is conducted by any statistical software is:

$$H_0 : \beta_k = 0$$
$$H_A : \beta_k \neq 0$$

Theory-Based Inference: Hypothesis Testing

- As we already mentioned, the default test that is conducted by any statistical software is:

$$H_0 : \beta_k = 0$$
$$H_A : \beta_k \neq 0$$

- So, **under the null hypothesis** we get from theory that the sampling distribution of $\frac{b}{\hat{\text{SE}}(b)}$ will be a standard normal distribution.

Theory-Based Inference: Hypothesis Testing

- As we already mentioned, the default test that is conducted by any statistical software is:

$$H_0 : \beta_k = 0$$
$$H_A : \beta_k \neq 0$$

- So, **under the null hypothesis** we get from theory that the sampling distribution of $\frac{b}{\hat{\text{SE}}(b)}$ will be a standard normal distribution.
- As such we can directly compare the observed test statistic $\frac{b}{\hat{\text{SE}}(b)}$ to the *standard normal distribution* which is the **null distribution** of our test statistic.

Theory-Based Inference: Hypothesis Testing

- As we already mentioned, the default test that is conducted by any statistical software is:

$$H_0 : \beta_k = 0$$
$$H_A : \beta_k \neq 0$$

- So, **under the null hypothesis** we get from theory that the sampling distribution of $\frac{b}{\hat{\text{SE}}(b)}$ will be a standard normal distribution.
- As such we can directly compare the observed test statistic $\frac{b}{\hat{\text{SE}}(b)}$ to the *standard normal distribution* which is the **null distribution** of our test statistic.
- The ***p-value*** associated to our test is then equal to the area of the *standard normal distribution* outside $\pm$ the observed value of $\frac{b}{\hat{\text{SE}}(b)}$.

Theory-Based Inference: Hypothesis Testing

- As we already mentioned, the default test that is conducted by any statistical software is:

$$H_0 : \beta_k = 0$$
$$H_A : \beta_k \neq 0$$

- So, **under the null hypothesis** we get from theory that the sampling distribution of $\frac{b}{\hat{\text{SE}}(b)}$ will be a standard normal distribution.
- As such we can directly compare the observed test statistic $\frac{b}{\hat{\text{SE}}(b)}$ to the *standard normal distribution* which is the **null distribution** of our test statistic.
- The ***p-value*** associated to our test is then equal to the area of the *standard normal distribution* outside $\pm$ the observed value of $\frac{b}{\hat{\text{SE}}(b)}$.
- Common rule of thumb: if the *estimate* is ***twice the size of the standard error***, then it is significant at the 5% level. Why?

Classical Regression Model

Classical Regression Model

- Whether the inference is made from theory or simulations, some assumptions have to be met for this inference to be valid.
- The set of assumptions needed defines the *Classical Regression Model* (CRM).

Classical Regression Model

- Whether the inference is made from theory or simulations, some assumptions have to be met for this inference to be valid.
- The set of assumptions needed defines the *Classical Regression Model* (CRM).
- Before delving into these assumptions, let's see the small but important modifications we apply to our model (back to *lecture 4: Simple Linear Regression*):

Classical Regression Model

- Whether the inference is made from theory or simulations, some assumptions have to be met for this inference to be valid.
- The set of assumptions needed defines the *Classical Regression Model* (CRM).
- Before delving into these assumptions, let's see the small but important modifications we apply to our model (back to *lecture 4: Simple Linear Regression*):
- We already mentioned the distinction between the sample estimate b_k (or $\hat{\beta}_k$) and the population parameter β_k.

Classical Regression Model

- Whether the inference is made from theory or simulations, some assumptions have to be met for this inference to be valid.
- The set of assumptions needed defines the *Classical Regression Model* (CRM).
- Before delving into these assumptions, let's see the small but important modifications we apply to our model (back to *lecture 4: Simple Linear Regression*):
- We already mentioned the distinction between the sample estimate b_k (or $\hat{\beta}_k$) and the population parameter β_k.
- In the same way, we distinguish e, the sample error, from ε the error term from the true population model:

$$y_i = \beta_0 + \beta_1 x_{1,i} + \ldots + \beta_k x_{k,i} + \varepsilon_i$$

CRM Assumptions

1. ***No perfect collinearity:*** the data are **not linearly dependent**, that is each variable provides new information on the outcome, and it is not a linear combination of the other variables.

CRM Assumptions

1. ***No perfect collinearity:*** the data are **not linearly dependent**, that is each variable provides new information on the outcome, and it is not a linear combination of the other variables.

2. ***Mean Independence:*** the mean of the residuals conditional on x should be zero, $E[\varepsilon|x] = 0$. Notice that this also means that $Cov(\varepsilon, x) = 0$, i.e. that the errors and our explanatory variable(s) should be *uncorrelated*.

CRM Assumptions

1. ***No perfect collinearity:*** the data are **not linearly dependent**, that is each variable provides new information on the outcome, and it is not a linear combination of the other variables.

2. ***Mean Independence:*** the mean of the residuals conditional on x should be zero, $E[\varepsilon|x] = 0$. Notice that this also means that $Cov(\varepsilon, x) = 0$, i.e. that the errors and our explanatory variable(s) should be *uncorrelated*.

3. ***Independently and identically distributed:*** the data are drawn from a **random sample** of size n: observation (x_i, y_i) comes from the exact same distribution, and is independent of observation (x_j, y_j), for all $i \neq j$.

CRM Assumptions

1. ***No perfect collinearity:*** the data are **not linearly dependent**, that is each variable provides new information on the outcome, and it is not a linear combination of the other variables.

2. ***Mean Independence:*** the mean of the residuals conditional on x should be zero, $E[\varepsilon|x] = 0$. Notice that this also means that $Cov(\varepsilon, x) = 0$, i.e. that the errors and our explanatory variable(s) should be *uncorrelated*.

3. ***Independently and identically distributed:*** the data are drawn from a **random sample** of size n: observation (x_i, y_i) comes from the exact same distribution, and is independent of observation (x_j, y_j), for all $i \neq j$.

4. ***Homoskedasticity:*** the variance of the error term ε is the same for each value of x: $Var(\varepsilon|x) = \sigma^2$.

CRM Assumptions

1. ***No perfect collinearity:*** the data are **not linearly dependent**, that is each variable provides new information on the outcome, and it is not a linear combination of the other variables.

2. ***Mean Independence:*** the mean of the residuals conditional on x should be zero, $E[\varepsilon|x] = 0$. Notice that this also means that $Cov(\varepsilon, x) = 0$, i.e. that the errors and our explanatory variable(s) should be *uncorrelated*.

3. ***Independently and identically distributed:*** the data are drawn from a **random sample** of size n: observation (x_i, y_i) comes from the exact same distribution, and is independent of observation (x_j, y_j), for all $i \neq j$.

4. ***Homoskedasticity:*** the variance of the error term ε is the same for each value of x: $Var(\varepsilon|x) = \sigma^2$.

5. ***Normally distributed errors:*** the error term is normally distributed, i.e. $\varepsilon \sim \mathcal{N}(0, \sigma^2)$

    - This last assumption allows avoiding large sample approximations, but it is never used in practice since samples are sufficiently large $(n \geq 30)$.

Exogeneity Assumption

The CRM assumption #2 is also know as the (strict) **exogeneity assumption**.

- When this assumption is violated our estimate b will be a ***biased*** estimate of β, i.e. $\mathbb{E}[b] \neq \beta$

Exogeneity Assumption

The CRM assumption #2 is also know as the (strict) **exogeneity assumption**.

- When this assumption is violated our estimate b will be a ***biased*** estimate of β, i.e. $\mathbb{E}[b] \neq \beta$
- For example, imagine you are interested in the effect of education on wage

$$\text{wage}_i = \beta_0 + \beta_1 \text{educ}_i + \varepsilon_i$$

  - Under the exogeneity assumption β_1 denotes the causal effect of education in the population.

Exogeneity Assumption

The CRM assumption #2 is also know as the (strict) **exogeneity assumption**.

- When this assumption is violated our estimate b will be a ***biased*** estimate of β, i.e. $\mathbb{E}[b] \neq \beta$
- For example, imagine you are interested in the effect of education on wage

$$\text{wage}_i = \beta_0 + \beta_1 \text{educ}_i + \varepsilon_i$$

  - Under the exogeneity assumption β_1 denotes the causal effect of education in the population.

- Suppose there is *unobserved* ability a_i.

  - High ability means higher wage.
  - It *also* means school is easier, and so i **selects** into more schooling.

Exogeneity Assumption

- Given ability is *unobserved*, a_i goes into the error ε_i

Exogeneity Assumption

- Given ability is *unobserved*, a_i goes into the error ε_i
- Our *ceteris paribus* assumption (all else equal) does not hold.

Exogeneity Assumption

- Given ability is *unobserved*, a_i goes into the error ε_i
- Our *ceteris paribus* assumption (all else equal) does not hold.
- Then regressing the wage on education we will attribute to `educ` part of the effect on wages that is actually *caused* by ability a_i!

Exogeneity Assumption

- Given ability is *unobserved*, a_i goes into the error ε_i
- Our *ceteris paribus* assumption (all else equal) does not hold.
- Then regressing the wage on education we will attribute to `educ` part of the effect on wages that is actually *caused* by ability a_i!
  - Remember the formula of the **omitted variable bias**:

$$\begin{aligned}\text{OVB} = &\{\text{Relationship between } ability_i \text{ and } educ_i\} \\ &* \{\text{Effect of } ability_i \text{ in multiple regression}\}\end{aligned}$$

Exogeneity Assumption

- Given ability is *unobserved*, a_i goes into the error ε_i
- Our *ceteris paribus* assumption (all else equal) does not hold.
- Then regressing the wage on education we will attribute to `educ` part of the effect on wages that is actually *caused* by ability a_i!
  - Remember the formula of the **omitted variable bias**:

$$\begin{aligned}\text{OVB} = &\{\text{Relationship between } ability_i \text{ and } educ_i\} \\ &* \{\text{Effect of } ability_i \text{ in multiple regression}\}\end{aligned}$$

  - Thus, we have:

$$\mathbb{E}(b_1) = \beta_1 + OVB > \beta_1$$

- *Interpretation*: taking repeated sample from the population and computing b_1 each time, we would **systematically overestimate** the effect of education on wage.

Breaking the other assumptions

- You can find examples associated to the other assumptions in our book!
- Takeaway: if assumptions violated, inference is invalid!

Task 3.1 (10 min)

Let's go back to our question of returns to education and gender.

1. Load the data `CPS1985` from the `AER` package and look back at the `help` to get the definition of each variable: `?CPS1985`. Call the data.frame `cps`.

2. Create the `log_wage` variable equal to the log of `wage`.

3. Regress `log_wage` on `gender` and `education`, and save it as `reg1`.

    - Interpret each coefficient.

    - Are the coefficients statistically significant? At which significance level?

4. Regress the `log_wage` on `gender`, `education` and their interaction `gender*education`, save it as `reg2`.

    - How do you interpret the coefficient associated to $female * education$?

    - Can we reject the nullity of this coefficient at the 5% level? At 10%?

Task 3.2 (10 min)

1. Produce a scatterplot of the relationship between the log wage and the level of education.

2. Add the *regression line* with `geom_smooth`. What does this line represents?

3. Let's illustrate what the shaded area stands for.

    1. Draw one bootstrap sample from our `cps` data.

    2. Regress the `log_wage` on `gender`, `education` and their interaction `gender*education`, save it as `reg_bootstrap`.

    3. From `reg_bootstrap` extract and save the value of the intercept for men as `intercept_men_bootstrap` and the value of the slope for men as `slope_men_bootstrap`. Do the same for women.

    4. Add both predicted lines from this bootstrap sample to the previous plot (*Hint*: use `geom_abline` (x2))

Illustrating Uncertainty

Let's repeat the procedure you just made 100 times!

```
library(AER)
data("CPS1985")
cps = CPS1985 %>% mutate(log_wage = log(wage))

set.seed(1)
bootstrap_sample = cps %>%
    rep_sample_n(size = nrow(cps), reps = 100, replace

ggplot(data=cps,aes(y = log_wage, x = education, colou
  geom_point(size = 1, alpha = 0.7) +
  geom_smooth(method = "lm", alpha = 2) +
  geom_smooth(data=bootstrap_sample,
              size = 0.2,
              aes(y = log_wage, x = education, group =
              method = "lm", se = FALSE) +
  facet_wrap(~gender) +
  scale_colour_manual(values = c("darkblue", "darkred"
  labs(x = "Education", y = "Log wage") +
  guides(colour=FALSE) +
  theme_bw(base_size = 20)
```

Illustrating Uncertainty

Let's repeat the procedure you just made 100 times!

```
library(AER)
data("CPS1985")
cps = CPS1985 %>% mutate(log_wage = log(wage))

set.seed(1)
bootstrap_sample = cps %>%
    rep_sample_n(size = nrow(cps), reps = 100, replace

ggplot(data=cps,aes(y = log_wage, x = education, colou
  geom_point(size = 1, alpha = 0.7) +
  geom_smooth(method = "lm", alpha = 2) +
  geom_smooth(data=bootstrap_sample,
              size = 0.2,
              aes(y = log_wage, x = education, group =
              method = "lm", se = FALSE) +
  facet_wrap(~gender) +
  scale_colour_manual(values = c("darkblue", "darkred"
  labs(x = "Education", y = "Log wage") +
  guides(colour=FALSE) +
  theme_bw(base_size = 20)
```

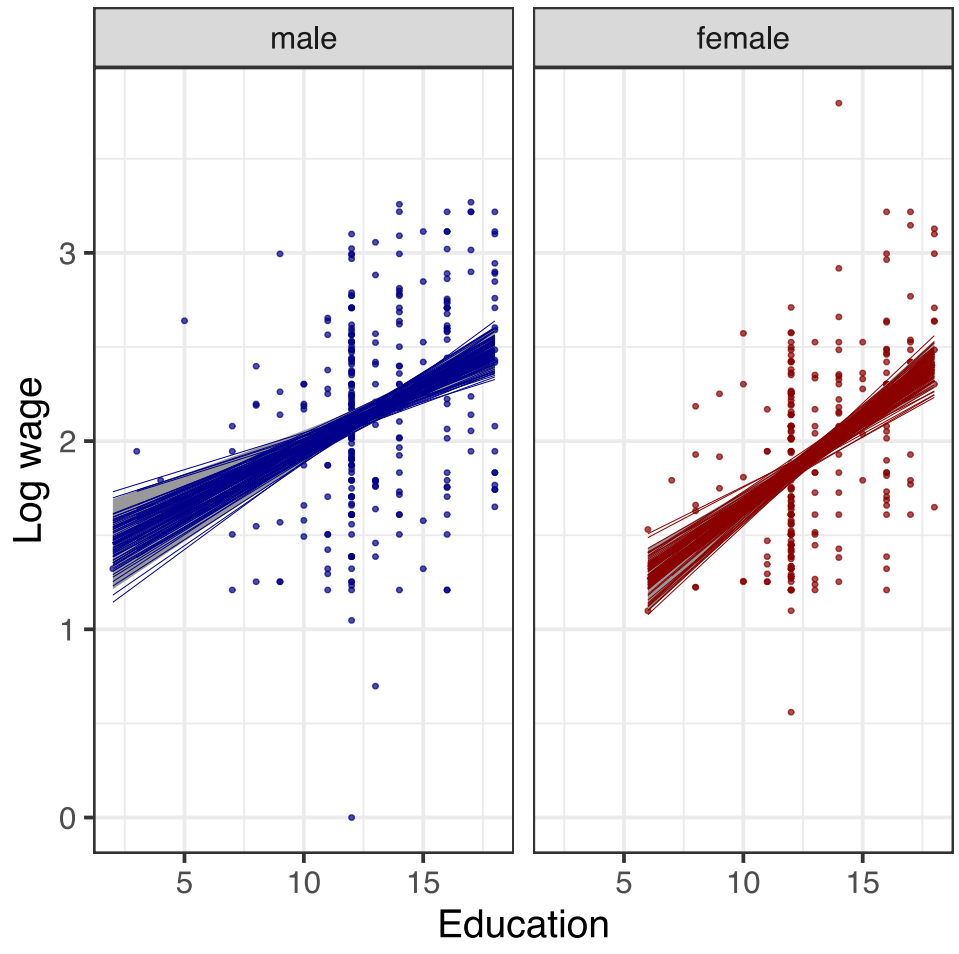

Illustrating Uncertainty

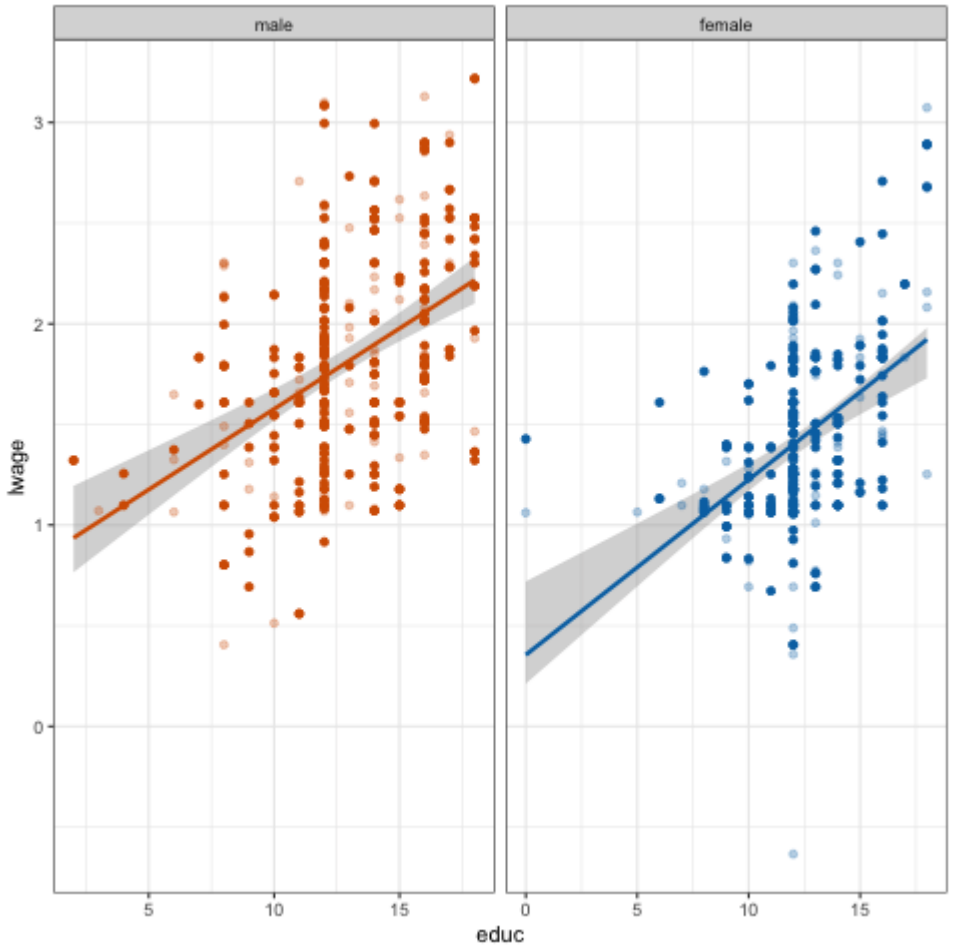

Even better : `ungeviz` and `gganimate` bring you moving lines!

- We took 20 bootstrap samples from our data
- You can see how different data points are included in each bootstrap sample.
- Those different points imply different regression lines.
- On average, 95% of these lines should fall into the shaded area.
- You should remember those moving lines when looking at the shaded area!

Teaser for next session

- Methods for program evaluation!

Teaser for next session

- Methods for program evaluation!
- 2 options:

Teaser for next session

- Methods for program evaluation!
- 2 options:

    1. Cover ***regression discontinuity design*** in-depth,

    OR

Teaser for next session

- Methods for program evaluation!
- 2 options:

  1. Cover ***regression discontinuity design*** in-depth,

  OR

  1. Cover both ***regression discontinuity design*** and ***differences-in-differences*** but you will only get the gist of each method.

Teaser for next session

- Methods for program evaluation!
- 2 options:

    1. Cover ***regression discontinuity design*** in-depth,

    OR

    1. Cover both ***regression discontinuity design*** and ***differences-in-differences*** but you will only get the gist of each method.

Which do you prefer?

THANKS

To the amazing moderndive team!

Big Thanks 🎉 to ungeviz and 🎊 gganimate for their awesome packages!

SciencesPo
DEPARTMENT OF ECONOMICS

SEE YOU NEXT WEEK!

	michele.fioretti@sciencespo.fr
	Slides
	Book
	@ScPoEcon
	@ScPoEcon